



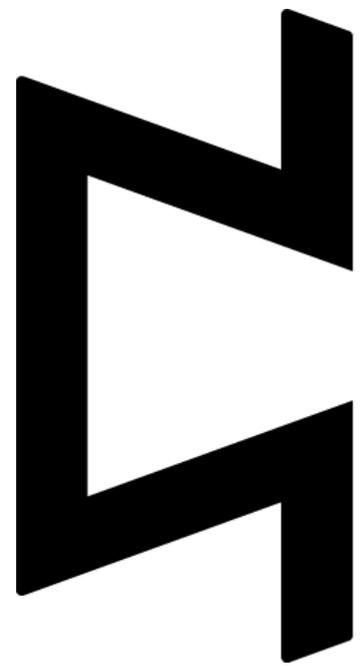
# Datomic

Copyright Cognitect, Inc.

This presentation is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License.

See <http://creativecommons.org/licenses/by-nc-sa/4.0/>

A Day of Datomic



cognitect

@stuarthalloway

# agenda

what is Datomic?

information model

transaction model

query model

time model

operational model

# what is Datomic?

a complete rethink of databases

agile

robust

powerful

time-aware

cloud

# agile: universal schema

entity / attribute / value / tx / op

granular, attribute-level schema definition

easy to model

easy to migrate

Data Set: [e a v t added], 257 items

e	a	v	t	added
:country/AF	:country/name	Afghanistan	13194139534626	true
:country/AL	:country/name	Albania	13194139534366	true
:country/DZ	:country/name	Algeria	13194139534392	true
:country/AS	:country/name	American Samoa	13194139534366	true
:country/AD	:country/name	Andorra	13194139534626	true

# update in place

sharing	difficult
distribution	difficult
concurrent access	difficult
access pattern	eager
caching	difficult
examples	Java and .NET collections relational databases NoSQL databases

**NOT  
ROBUST**

# persistent data structures

sharing	trivial
distribution	easy
concurrent access	trivial
access pattern	eager or lazy
caching	easy
examples	Clojure, F# collections Datomic database

**ROBUST**

# powerful

universal schema supports many access styles

row

column, key-value, document, and graph

declarative

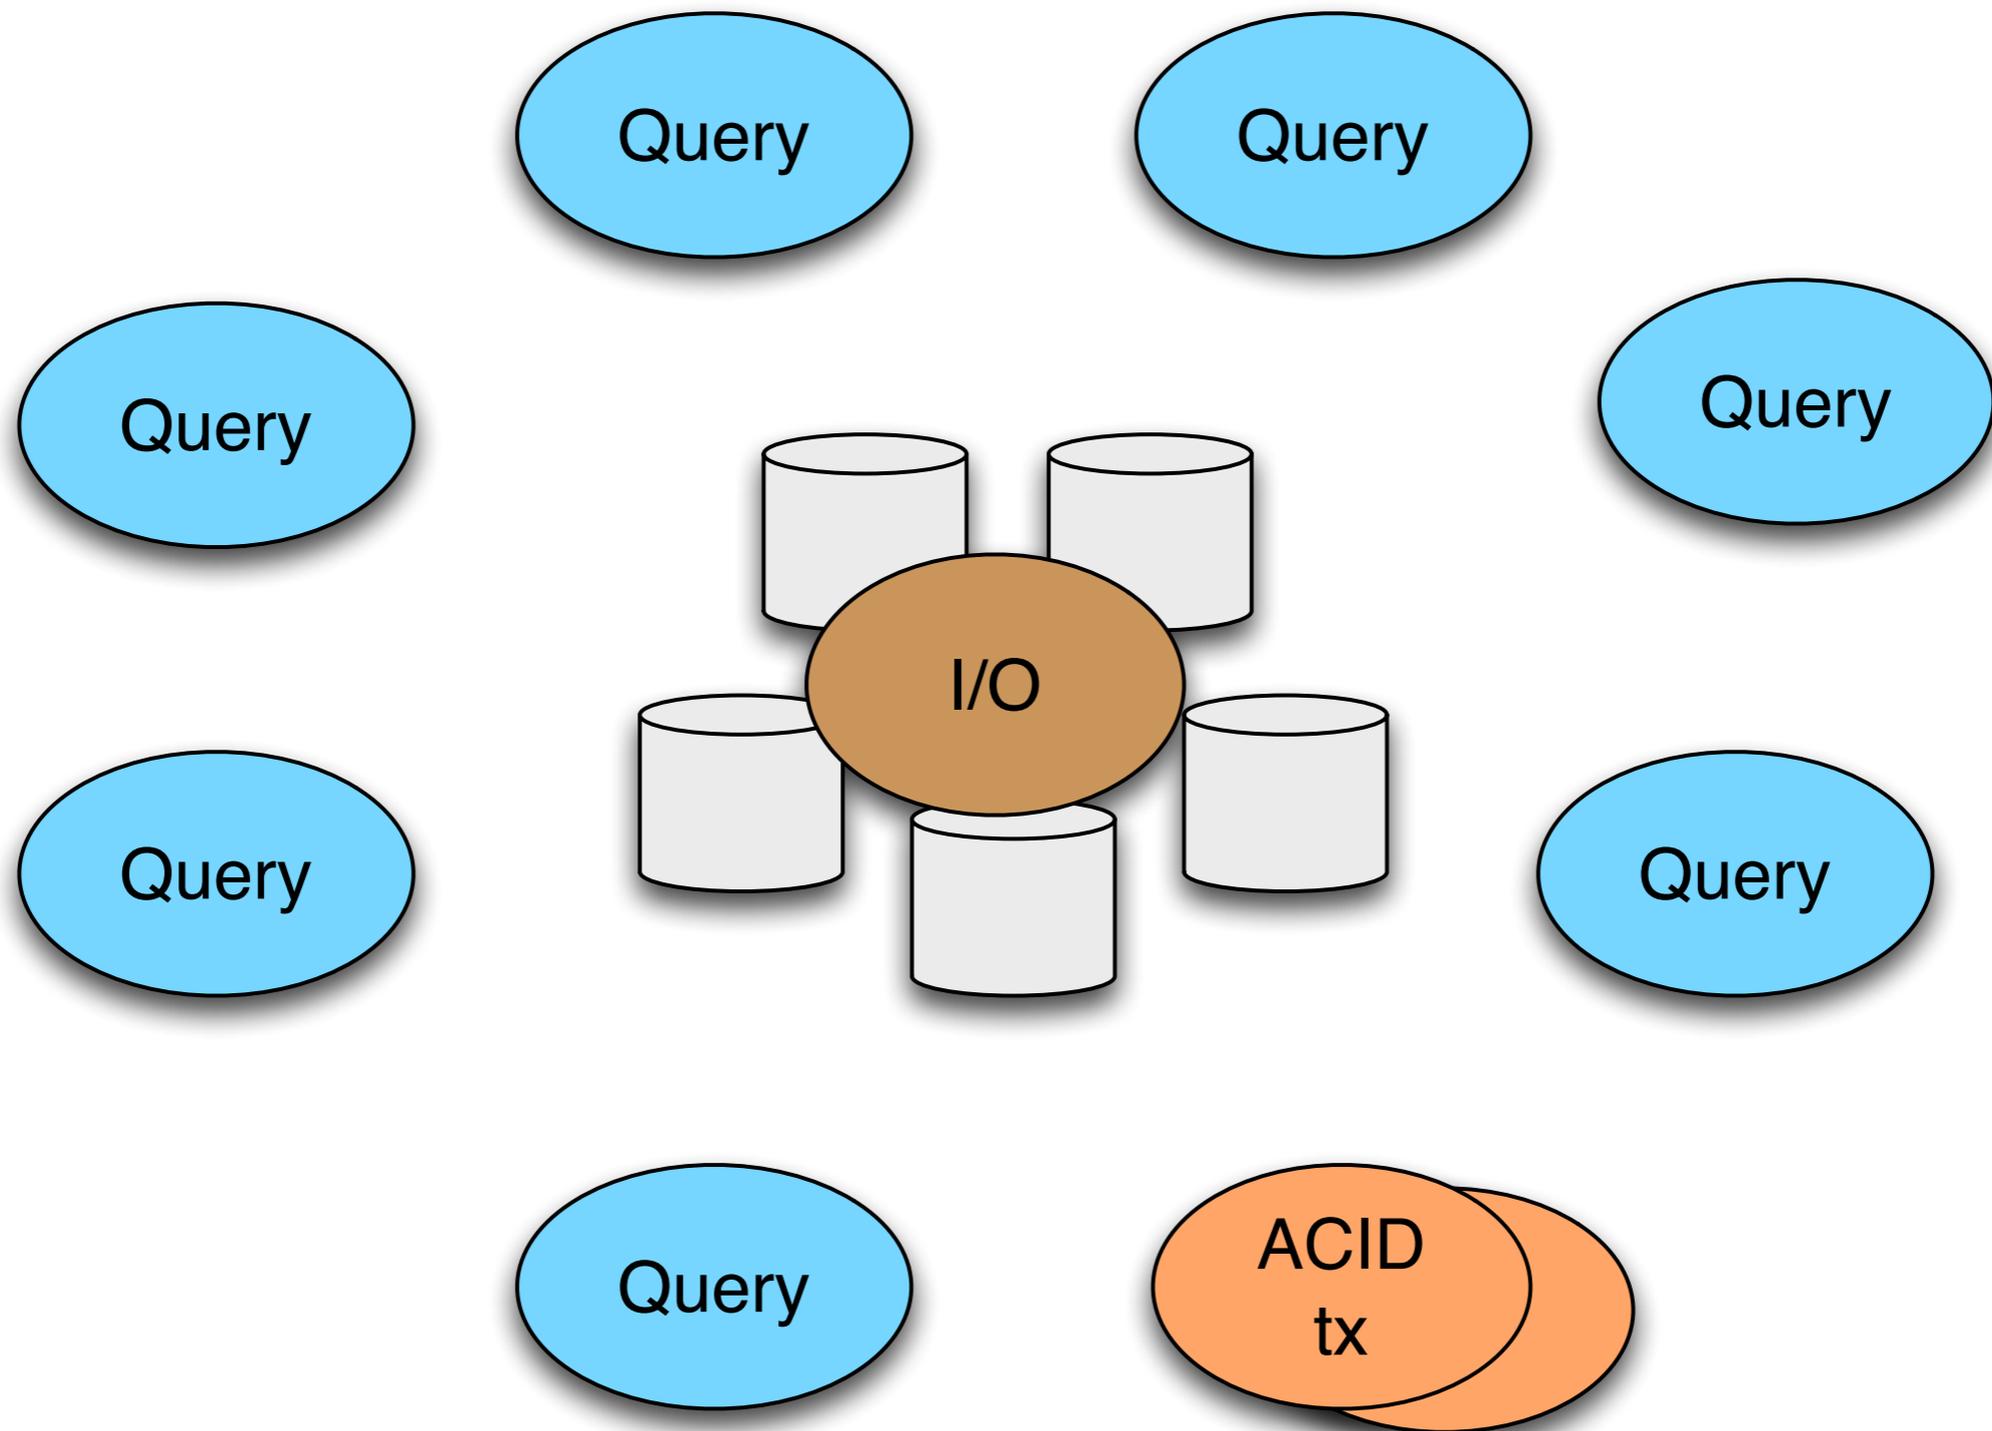
datalog query

pull

# time aware

db view	semantics	supports
<i>(default)</i>	current state	what is the current situation?
.asOf	state at point in past	how were things in the past?
.since	state since point in past	how have things changed?
tx report	before / after / change view of tx	automated event response
.with	state with proposed additions	what would happen if we did X?
.history	timeless view of all history	anything!

# a database of little services



# lab: running the tools

install Datomic, start dev transactor	<a href="http://docs.datomic.com/getting-started.html">http://docs.datomic.com/getting-started.html</a>
install mbrainz dataset to datomic:dev://localhost:4334/mbrainz	<a href="https://github.com/Datomic/mbrainz-sample">https://github.com/Datomic/mbrainz-sample</a>
explore mbrainz from console	<a href="http://docs.datomic.com/console.html">http://docs.datomic.com/console.html</a>
explore mbrainz from REST service	<a href="http://docs.datomic.com/rest.html">http://docs.datomic.com/rest.html</a>
install examples	<a href="https://github.com/datomic/day-of-datomic">https://github.com/datomic/day-of-datomic</a>
install your favorite REPL (links show Cursive)	<a href="https://www.jetbrains.com/idea/download/">https://www.jetbrains.com/idea/download/</a> + <a href="https://cursiveclojure.com/userguide/">https://cursiveclojure.com/userguide/</a>
play through CRUD example	<a href="https://github.com/Datomic/day-of-datomic/blob/master/tutorial/crud.clj">https://github.com/Datomic/day-of-datomic/blob/master/tutorial/crud.clj</a>

# information model

notation

datoms

databases

entities

schema

# notation: edn

```
{ :firstName "John"  
  :lastName "Smith"  
  :age 25  
  :address {  
    :streetAddress "21 2nd Street"  
    :city "New York"  
    :state "NY"  
    :postalCode "10021" }  
  :phoneNumber  
  [ { :type "name" :number "212 555-1234"}  
    { :type "fax" :number "646 555-4567" } ] }
```

type	examples
string	<code>"foo"</code>
character	<code>\f</code>
integer	<code>42, 42N</code>
floating point	<code>3.14, 3.14M</code>
boolean	<code>true</code>
nil	<code>nil</code>
symbol	<code>foo, +</code>
keyword	<code>:foo, ::foo</code>

type	properties	examples
list	sequential	<code>(1 2 3)</code>
vector	sequential and random access	<code>[1 2 3]</code>
map	associative	<code>{:a 100 :b 90}</code>
set	membership	<code>#{:a :b}</code>

# generic extensibility

*#name edn-form*

name describes interpretation of following element

recursively defined

all data can be literal

# built-in tags

**#inst "rfc-3339-format"**

tagged element is a string in RFC-3339 format

**#uuid "f81d4fae-7dec-11d0-a765-00a0c91e6bf6"**

tagged element is a canonical UUID string

# datoms

granular, atomic facts

immutable

5-tuple: entity / attribute / value / transaction / op

# example datoms

e	a	v	tx	op
jane	likes	broccoli	1008	true
jane	likes	pizza	1008	true
jane	likes	pizza	1148	false

# datom API

```
public interface Datom {  
    Object e();  
    Object a();  
    Object v();  
    Object tx();  
    boolean added();  
  
    // positional  
    Object get(int index);  
}
```

# database

a set of datoms (universal relation)

efficient storage

many sort orders

accumulate-only (not append-only)

# database sorts in console

Query Entities Transactions **Indexes**

Index

A

V

E

T

Data Set: [e a v t added], 257 items

e	a	v	t	added
:country/AF	:country/name	Afghanistan	13194139534626	true
:country/AL	:country/name	Albania	13194139534366	true
:country/DZ	:country/name	Algeria	13194139534392	true
:country/AS	:country/name	American Samoa	13194139534366	true
:country/AD	:country/name	Andorra	13194139534626	true

# entities

immutable

lazy

associative

inferred from datoms sharing a common **e**

point-in-time

bidirectional navigation

# example entity

entity

```
{:db/id jane  
 :likes "broccoli"}
```



datoms

e	a	v	tx	op
jane	likes	broccoli	1008	true
jane	likes	pizza	1008	true
jane	likes	pizza	1148	false

# entity API

```
interface Database
{
    Entity entity(Object entityId);
    // unrelated methods elided for brevity
}

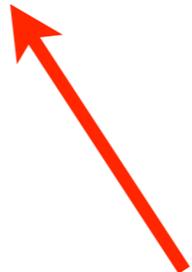
interface Entity {
    Object get(Object key);
    Entity touch();
    Set keySet();
}
```

# relation direction

```
// things that Jane likes  
jane.get("likes");
```

# reversing direction

```
// people that like broccoli  
broccoli.get("_likes");
```



underscore means  
“nav backwards”

# entities in console

The screenshot shows a console interface with a tabbed view. The 'Entities' tab is active, displaying the details for an entity with ID ':country/AL'. The details are organized into a table with two columns: 'Attribute' and 'Value'. The attributes shown are ':db/id' (value: 17592186045483), ':country/name' (value: Albania), and ':db/ident' (value: :country/AL). A red arrow points from the ':country/AL' text in the 'Entity ID' field to the ':country/AL' text in the 'Attribute' column of the table below.

Attribute	Value
:db/id	17592186045483
:country/name	Albania
:db/ident	:country/AL

Data Set: [e a v t added], 257 items

e	a	v	t	added
:country/AF	:country/name	Afghanistan	13194139534626	true
:country/AL	:country/name	Albania	13194139534366	true

click any e

# schema

*schemas add power*

schema is plain data

schema elements installed via transactions

make history-compatible changes at any time

# common schema attributes

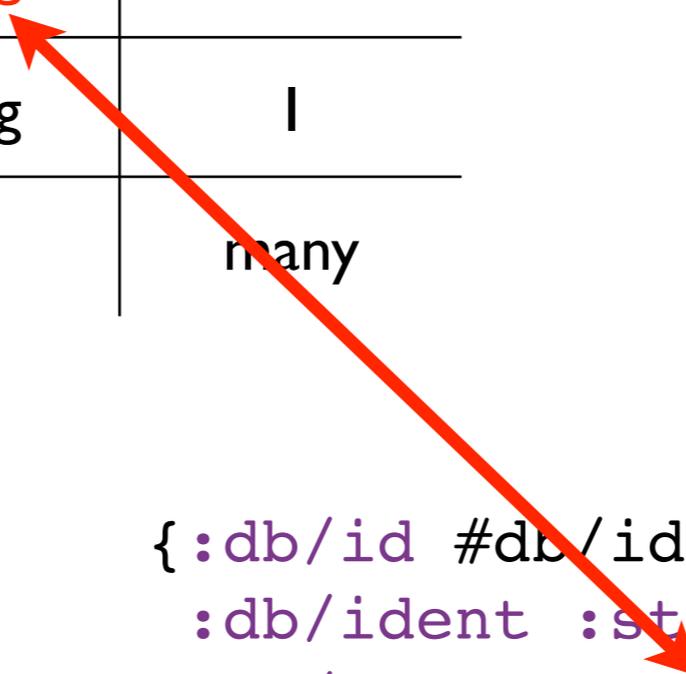
attribute	type	use
db/ident	keyword	programmatic name
db/valueType	ref	attribute type
db/cardinality	ref	one- or many- valued?
db/index	ref	creates AVET
db/unique	ref	unique, “upsert”
db/isComponent	ref	ownership

# stories

<b>attribute</b>	<b>type</b>	<b>cardinality</b>
story/title	string	1
story/url	string	1
story/slug	string	1
news/comments	ref	many

# schema is plain old data

attribute	type	card
story/title	string	1
story/url	string	1
story/slug	string	1
news/comments	ref	many



```
{:db/id #db/id[:db.part/db]  
:db/ident :story/url  
:db/valueType :db.type/string  
:db/cardinality :db.cardinality/one  
:db.install/_attribute :db.part/db}
```

# users

attribute	type	cardinality
user/firstName	string	1
user/lastName	string	1
user/email*	string	1
user/upVotes	ref	many

\*unique

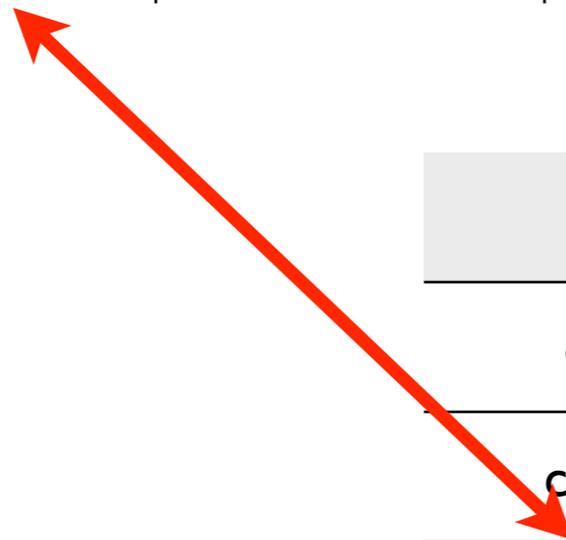
# comments

attribute	type	cardinality
comment/body	string	1
comment/author	ref	1
news/comments	ref	many

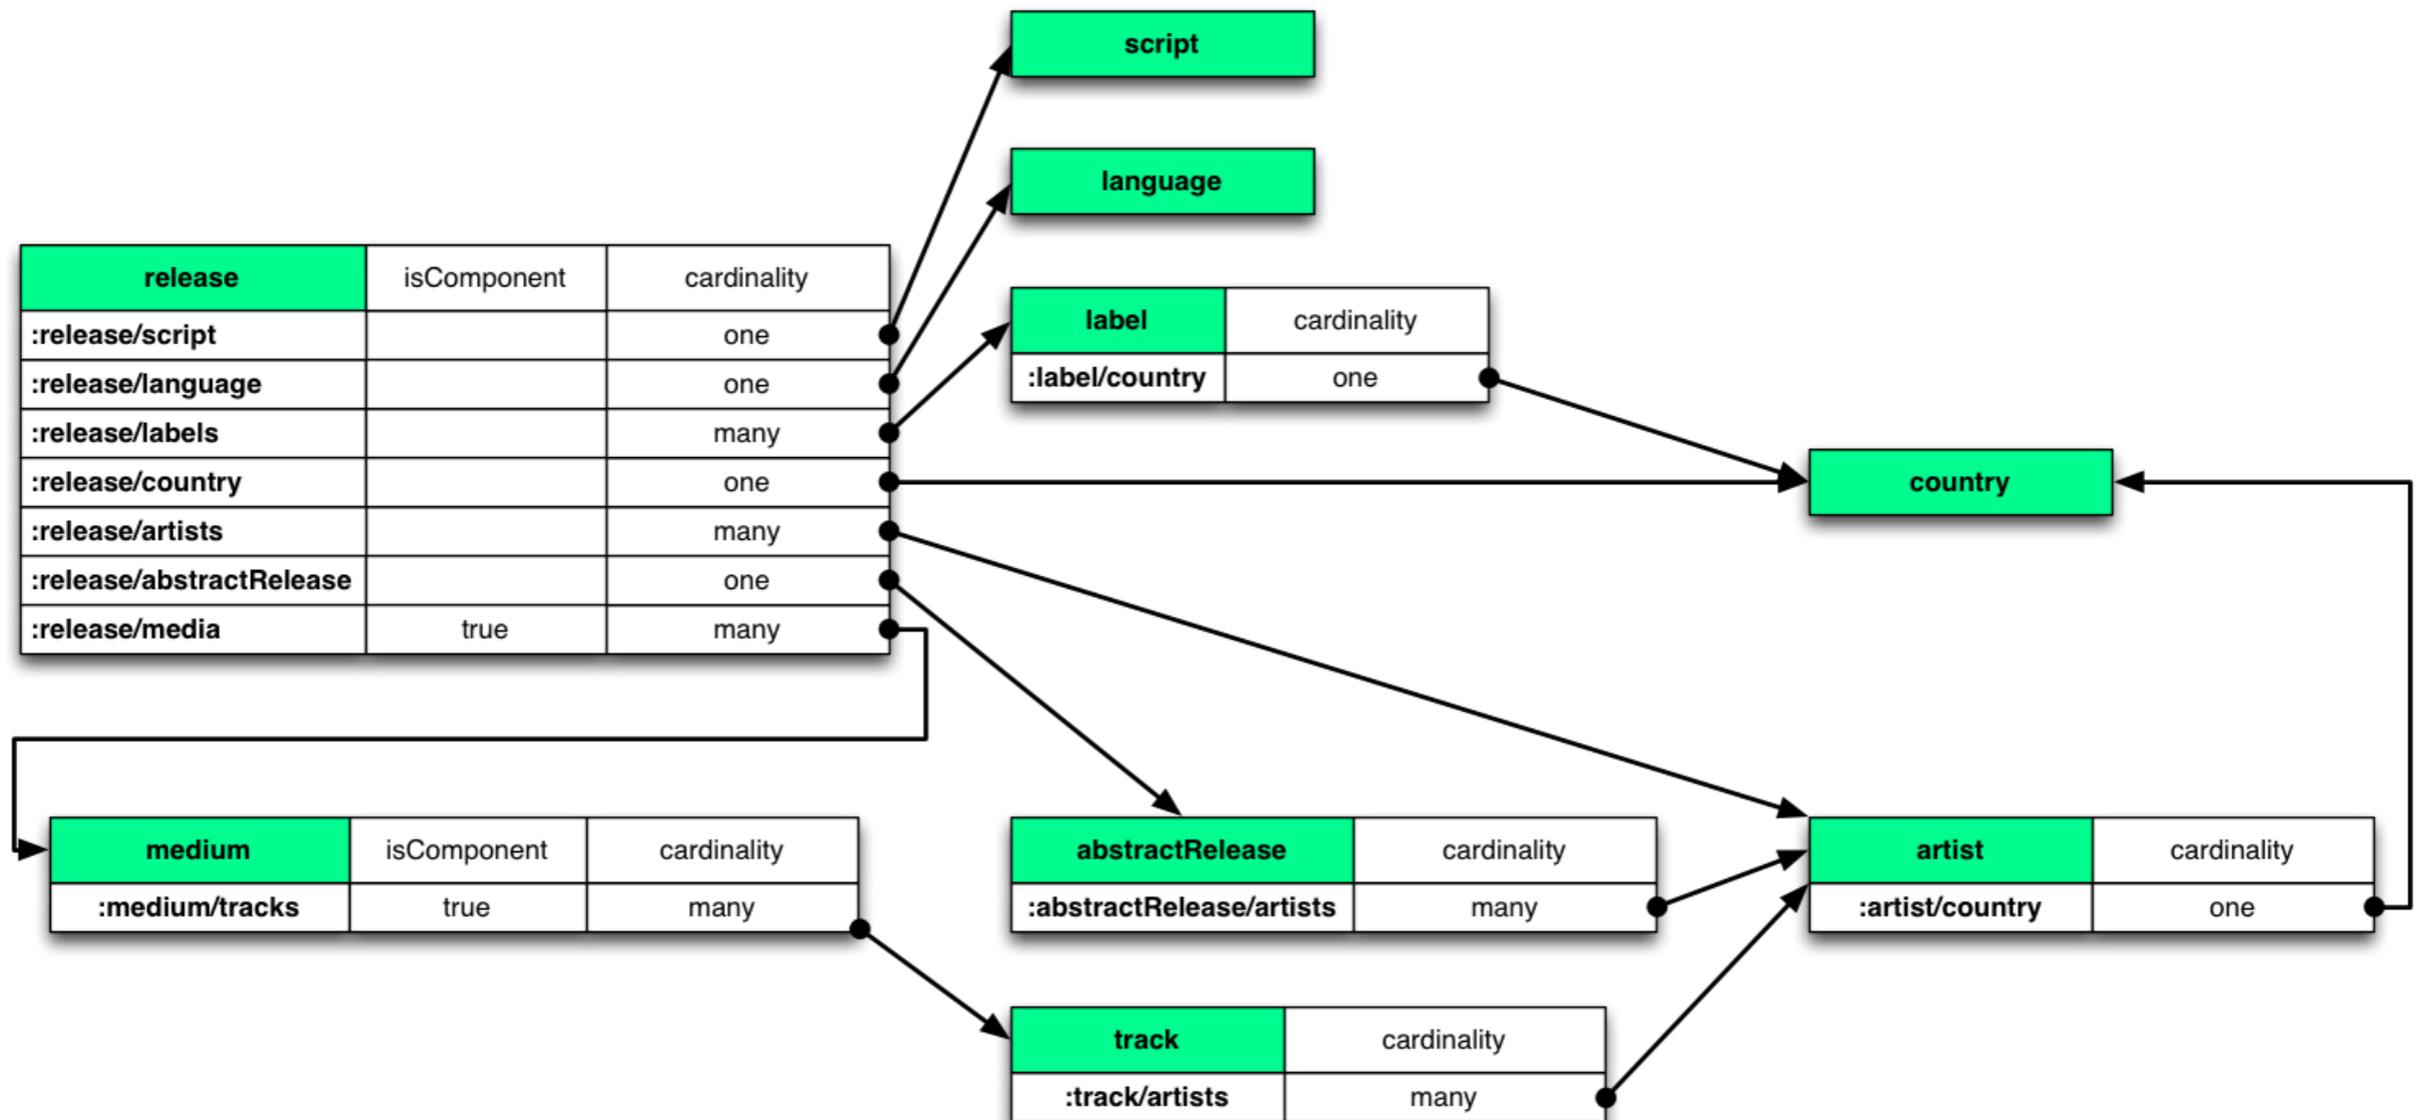
# “types” do not dictate attrs

attribute	type	cardinality
story/title	string	1
story/url	string	1
story/slug	string	1
news/comments	ref	many

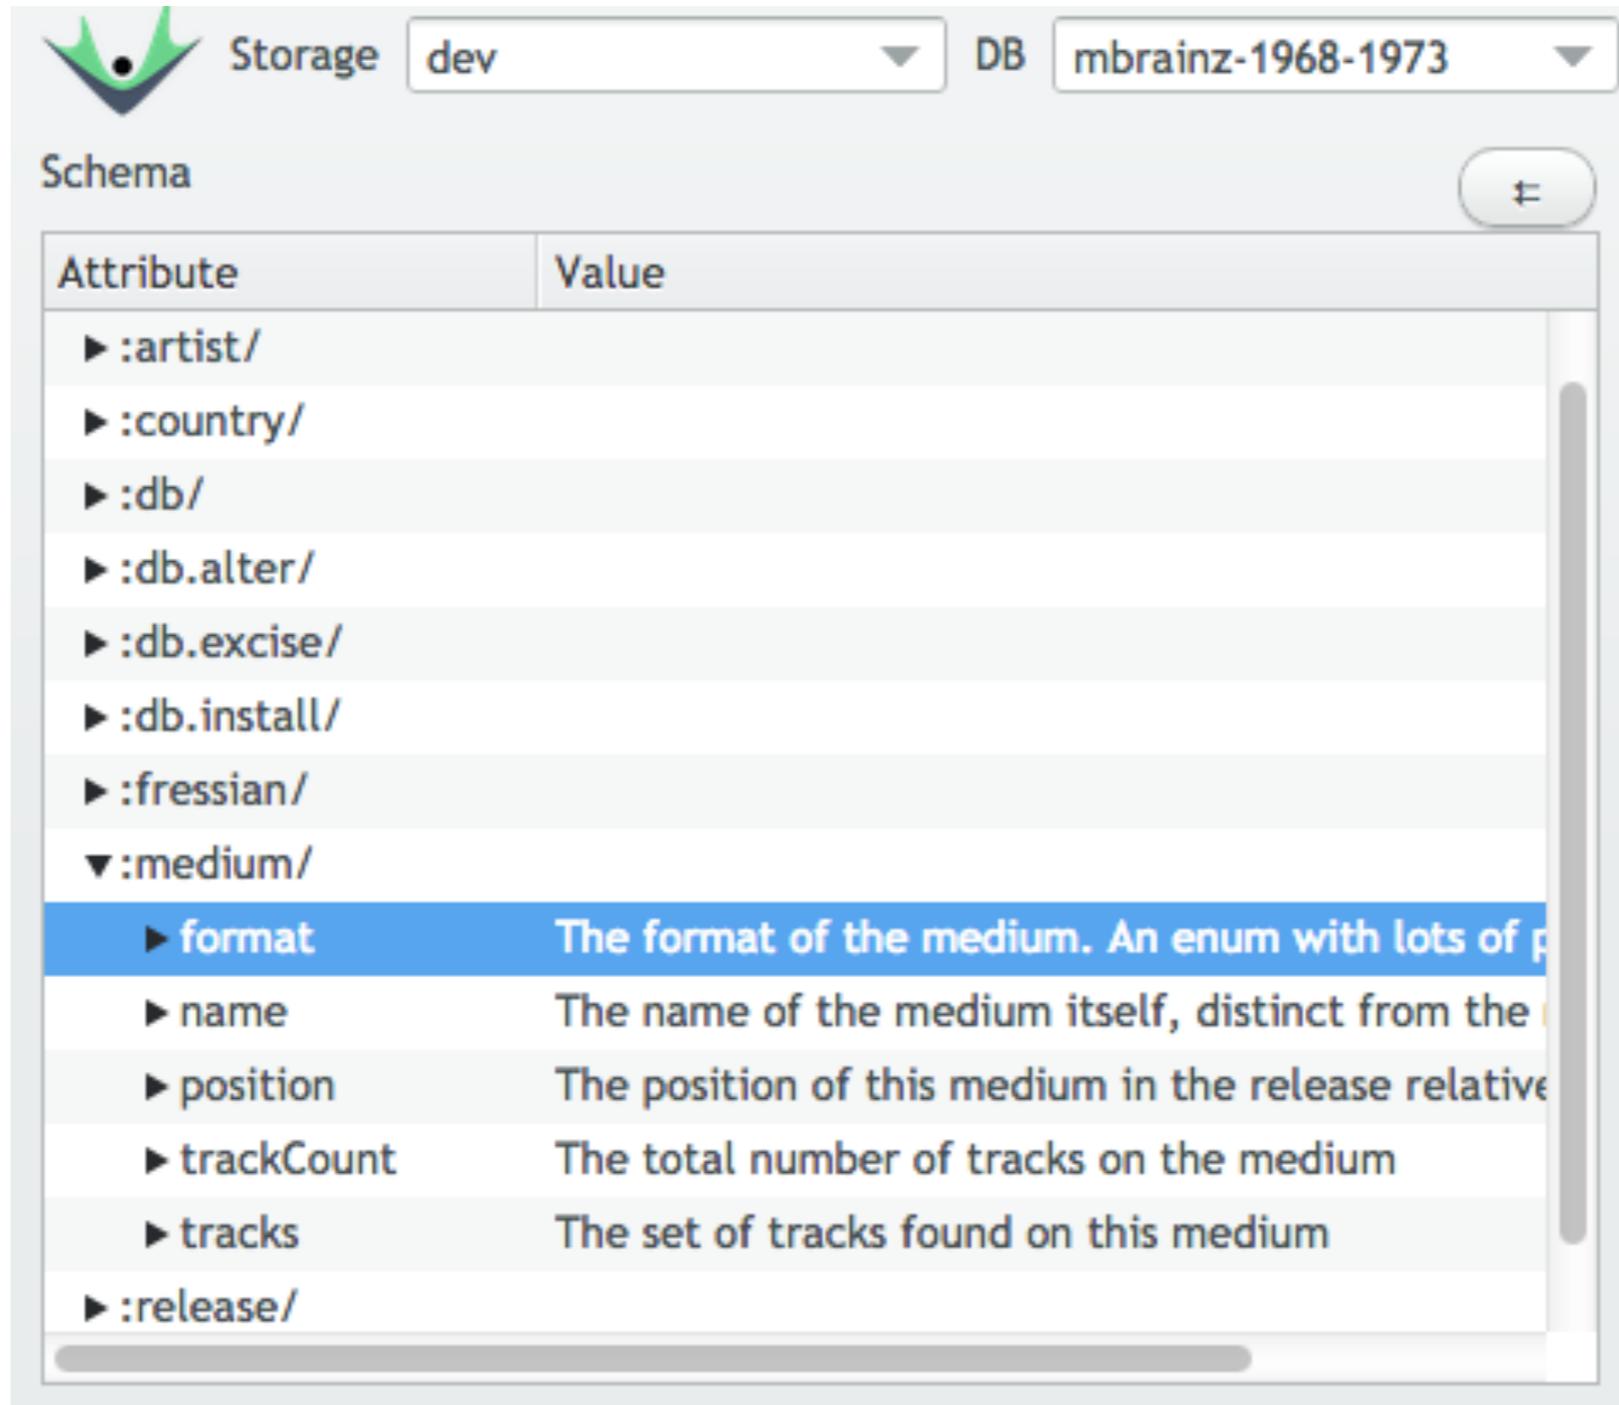
attribute	type	cardinality
comment/body	string	1
comment/author	ref	1
news/comments	ref	many



# example: mbrainz



# schema in console



Storage dev DB mbrainz-1968-1973

Schema

Attribute	Value
▶ :artist/	
▶ :country/	
▶ :db/	
▶ :db.alter/	
▶ :db.excise/	
▶ :db.install/	
▶ :fressian/	
▼ :medium/	
▶ format	The format of the medium. An enum with lots of p
▶ name	The name of the medium itself, distinct from the
▶ position	The position of this medium in the release relative
▶ trackCount	The total number of tracks on the medium
▶ tracks	The set of tracks found on this medium
▶ :release/	

# lab: designing a schema

pick a **small** domain you know well

draw an entity/relationship diagram

convert to edn data

transact into a database

view schema in console

# transaction model

ACID

assertion and retraction

entity maps

ids and partitions

uniqueness

transaction functions

# ACID

atomic	transaction is a set of datoms transaction entirely in single write
consistent	all processes see same global ordering of transactions
isolated	single writer system (nobody to be isolated <i>from</i> )
durable	always flush through to durable storage before reporting transaction complete

# assertion and retraction

```
[ :db/add entity-id attribute value]
```

```
[ :db/retract entity-id attribute value]
```

# entity maps

concise form for multiple assertions about an entity

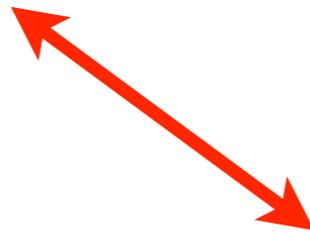
equivalent to corresponding list of asserts

can nest arbitrarily

# lists vs. entity maps

```
[ :db/add 42 :likes "pizza" ]  
[ :db/add 42 :firstName "John" ]  
[ :db/add 42 :lastName "Doe" ]
```

assertion  
list



```
{ :db/id 42  
  :likes "pizza"  
  :firstName "John"  
  :lastName "Doe" }
```

entity  
map

# cross reference

```
[ { :db/id #db/id[ :db.part/user -1 ]  
  :person/name "Bob"  
  :person/spouse #db/id[ :db.part/user -2 ] }
```

```
{ :db/id #db/id[ :db.part/user -2 ]  
  :person/name "Alice"  
  :person/spouse #db/id[ :db.part/user -1 ] } ]
```

# nesting

```
[ { :db/id order-id  
  :order/lineItems [ { :lineItem/product chocolate  
                      :lineItem/quantity 1 }  
                    { :lineItem/product whisky  
                      :lineItem/quantity 2 } ] ] }
```

# partitions

# partitions

group related entities in a partition

coarser granularity than e.g. tables

partition is a hint to indexing

group these things together

can help locality

does not affect semantics

# partition locality

<b>part</b>	<b>e</b>	<b>a</b>	<b>v</b>	<b>tx</b>	<b>added</b>
0	0x0000000000000000	10	:db.part/db	0xc000000000	TRUE
0	0x0000000000000000	11	0	0xc0000000036	TRUE
0	0x0000000000000000	11	3	0xc0000000036	TRUE
0	0x0000000000000000	11	4	0xc0000000036	TRUE
3	0x0000c00000000036	50	Wed Dec 31 19:00:00 E...	0xc0000000036	TRUE
3	0x0000c00000000038	50	Wed Dec 31 19:00:00 E...	0xc0000000038	TRUE
3	0x0000c0000000003f	50	Wed Dec 31 19:00:00 E...	0xc000000003f	TRUE
3	0x0000c000000003e8	50	Mon Oct 13 18:52:59 E...	0xc00000003e8	TRUE
4	0x00001000000003eb	10	:country/BF	0xc00000003ea	TRUE
4	0x00001000000003eb	84	Burkina Faso	0xc00000003ea	TRUE
4	0x00001000000003ec	10	:country/JE	0xc00000003ea	TRUE
4	0x00001000000003ec	84	Jersey	0xc00000003ea	TRUE

# built-in partitions

<b>partition</b>	<b>usage</b>
<code>:db.part/db</code>	schema entities
<code>:db.part/tx</code>	transaction entities
<code>:db.part/user</code>	user entities

# creating partitions

```
[ { :db.install/_partition :db.part/db,  
  :db/id #db/id[:db.part/db],  
  :db/ident :inventory }  
{ :db.install/_partition :db.part/db,  
  :db/id #db/id[:db.part/db],  
  :db/ident :customers } ]
```

# identity

# identity

requirement	model with	value types
db-relative opaque id	entity id	opaque (long)
external id	:db.unique/identity attribute	string, uuid, uri
global opaque id	:db.unique/identity squuid	uuid
programmatically name	:db/ident	keyword

# squids

semi-sequential UUIDs

do not fragment indexes

```
public class Peer;  
    public static UUID squuid();  
    public static long squuidTimeMillis(UUID squuid);  
    // other methods elided for brevity  
}
```

# transaction functions

# add and retract

```
[[:db/add john :likes pizza]  
[:db/retract john :likes iceCream]]
```

# what about update?

```
[[:db/add john :likes pizza]  
[:db/retract john :likes iceCream]  
[:db/add john :balance 110?]]
```

# atomic increment

```
[[:db/add john :likes pizza]  
[:db/retract john :likes iceCream]  
[:inc john :account 10]]
```

# transaction fns

subset of data fns

run inside transactions

have access to in-tx value of database

as first argument

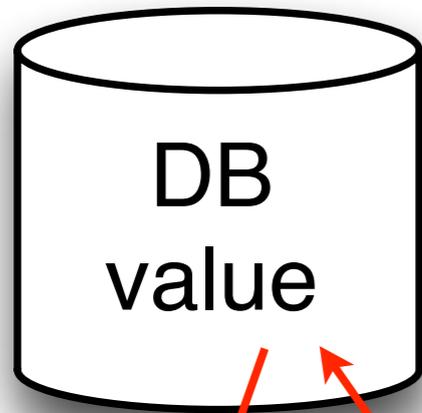
# tx function expansion

```
[[:db/add john :likes pizza]  
[:db/retract john :likes iceCream]  
[:inc john :balance 10]]
```



```
[[:db/add john :likes pizza]  
[:db/retract john :likes iceCream]  
[:db/add john :balance 110]]
```

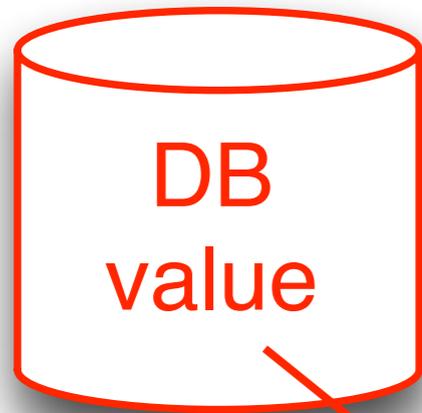
# lookup the function



```
[ :inc john :balance 10 ]
```

```
inc = db.entity("inc").get("db/fn");
```

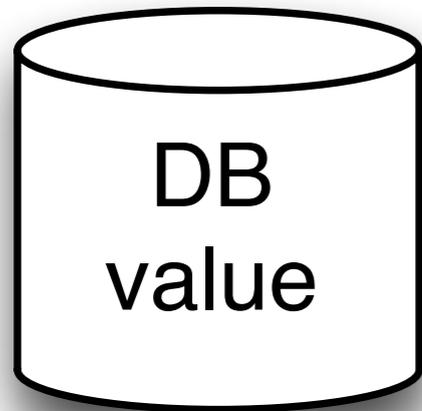
# pass in current db



```
[ :inc john :account 10 ]
```

```
inc.invoke(db, ...);
```

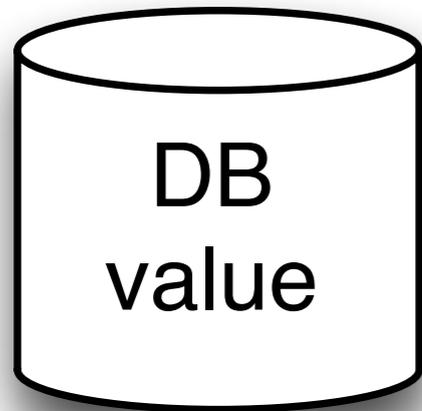
# pass in args



```
[ :inc john :account 10 ]
```

```
inc.invoke(db, john, :account, 10)
```

# data out



```
[ :inc john :account 10 ]
```

```
inc.invoke(db, john, :account, 10)
```



```
[ [ :db/add john :account 110 ] ]
```

# inc

```
public static Object inc(Object db, Object e, Object amount)
{
    // lookup entity
    // calculate new balance
    // create assertion
    // return list containing assertion
}
```

# inc

```
public static Object inc(Object db, Object e, Object a, Object amount) {  
    Entity ent = ((Database)db).entity(e);  
    Long balance = (Long) ent.get(a) + (Long) amount;  
    List updated = list("db/add", e, a, balance);  
    return list(updated);  
}
```

# lab: adding some data

create assertions and retractions	<a href="https://github.com/Datomic/day-of-datomic/blob/master/tutorial/crud.clj">https://github.com/Datomic/day-of-datomic/blob/master/tutorial/crud.clj</a>
create entity maps	<a href="https://github.com/Datomic/day-of-datomic/blob/master/tutorial/component_attributes.clj">https://github.com/Datomic/day-of-datomic/blob/master/tutorial/component_attributes.clj</a>
modify an existing entity	<a href="https://github.com/Datomic/day-of-datomic/blob/master/tutorial/crud.clj">https://github.com/Datomic/day-of-datomic/blob/master/tutorial/crud.clj</a>
review your data at the console	
bonus: create a constructor function	<a href="https://github.com/Datomic/day-of-datomic/blob/master/tutorial/data_functions.clj">https://github.com/Datomic/day-of-datomic/blob/master/tutorial/data_functions.clj</a>

# query model

datalog

pull

entities

raw indexes

# why datalog?

equivalent to relational model + recursion

better fit than prolog for query

- no clause order dependency

- guaranteed termination

pattern-matching style easy to learn

# example database

entity	attribute	value
42	:email	<u>jd</u> oe@example.com
43	:email	<u>jan</u> e@example.com
42	:orders	107
42	:orders	141

# data pattern

*Constrains the results returned,  
binds variables*

```
[?customer :email ?email]
```

# data pattern

*Constrains the results returned,  
binds variables*

[?customer :email ?email]



entity



attribute



value

# data pattern

*Constrains the results returned,  
binds variables*

constant



[?customer :email ?email]

# data pattern

*Constrains the results returned,  
binds variables*

variable



variable



[?customer :email ?email]

# example database

entity	attribute	value
42	:email	<u>jdoe@example.com</u>
43	:email	<u>jane@example.com</u>
42	:orders	107
42	:orders	141

[ ?customer :email ?email ]

# constants anywhere

“Find a particular customer’s email”

```
[ 42 :email ?email ]
```

entity	attribute	value
42	:email	<u>jdoue@example.com</u>
43	:email	<u>jane@example.com</u>
42	:orders	107
42	:orders	141

[ 42 :email ?email ]

# variables anywhere

“What attributes does  
customer 42 have?”

```
[ 42 ?attribute ]
```

entity	attribute	value
42	:email	<u>jdoe@example.com</u>
43	:email	<u>jane@example.com</u>
42	:orders	107
42	:orders	141

[ 42 ?attribute ]

# variables anywhere

“What attributes and values does customer 42 have?”

[ 42 **?attribute ?value** ]

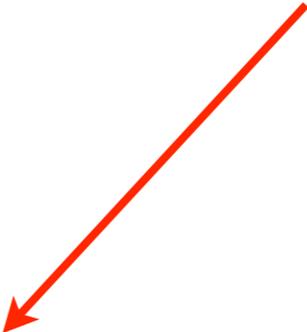
entity	attribute	value
42	:email	<u>jdoe@example.com</u>
43	:email	<u>jane@example.com</u>
42	:orders	107
42	:orders	141

[ 42 ?attribute ?value ]

# where clause

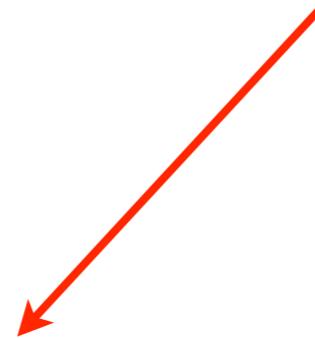
```
[ :find ?customer  
  :where [ ?customer :email ] ]
```

data  
pattern



# find clause

variable to  
return



```
[ :find ?customer  
  :where [?customer :email]]
```

# implicit join

“Find all the customers who have placed orders.”

```
[ :find ?customer  
  :where [ ?customer :email ]  
          [ ?customer :orders ] ]
```

# API

```
import static datomic.Peer.q;
```

```
q("[:find ?customer  
  :where [?customer :id]  
         [?customer :orders]]",  
  db);
```

# q

```
import static datomic.Peer.q;
```

```
q("[:find ?customer  
  :where [?customer :id]  
         [?customer :orders]]",  
  db);
```

# query

```
import static datomic.Peer.q;
```

```
q("[:find ?customer  
  :where [?customer :id]  
         [?customer :orders]]",  
  db);
```

# input(s)

```
import static datomic.Peer.q;
```

```
q("[:find ?customer  
  :where [?customer :id]  
         [?customer :orders]]",  
  db);
```

# in clause

*Names inputs so you can refer to them elsewhere in the query*

```
:in $database ?email
```

# parameterized query

“Find a customer by email.”

```
q([:find ?customer  
  :in $database ?email  
  :where [$database ?customer :email ?email]],  
db,  
"jdoe@example.com");
```

# first input

“Find a customer by email.”

```
q([:find ?customer  
  :in $database ?email  
  :where [$database ?customer :email ?email]],  
  db,  
  "jdoe@example.com" );
```

# second input

“Find a customer by email.”

```
q([:find ?customer  
  :in $database ?email  
  :where [$database ?customer :email ?email]],  
db,  
"jdoe@example.com");
```

# verbose?

“Find a customer by email.”

```
q([:find ?customer  
  :in $database ?email  
  :where [$database ?customer :email ?email]],  
  db,  
  "jdoe@example.com" );
```

# shortest name possible

“Find a customer by email.”

```
q([:find ?customer  
  :in $ ?email  
  :where [$ ?customer :email ?email]],  
  db,  
  "jdoe@example.com");
```

# elide \$ in where

“Find a customer by email.”

```
q([:find ?customer  
  :in $ ?email  
  :where [ ?customer :email ?email]],  
db,  
"jdoe@example.com");
```



no need to  
specify \$

# predicates

*Functional constraints that can  
appear in a :where clause*

```
[ (< 50 ?price) ]
```

# adding a predicate

“Find the expensive items”

```
[ :find ?item  
  :where [?item :item/price ?price]  
         [ (< 50 ?price) ] ]
```

# functions

*Take bound variables as inputs  
and bind variables with output*

```
[ ( shipping ?zip ?weight ) ?cost ]
```

# function args

```
[ ( shipping ?zip ?weight ) ?cost ]
```



bound inputs

# function returns

```
[ (shipping ?zip ?weight) ?cost ]
```

bind return  
values



# calling a function

“Find me the customer/product combinations where the shipping cost dominates the product cost.”

```
[ :find ?customer ?product
  :where [?customer :shipAddress ?addr]
         [?addr :zip ?zip]
         [?product :product/weight ?weight]
         [?product :product/price ?price]
         [(Shipping/estimate ?zip ?weight) ?shipCost]
         [(<= ?price ?shipCost)]]
```

← navigate from customer to zip

# calling a function

“Find me the customer/product combinations where the shipping cost dominates the product cost.”

```
[ :find ?customer ?product
  :where [?customer :shipAddress ?addr]
         [?addr :zip ?zip]
         [?product :product/weight ?weight]
         [?product :product/price ?price]
         [(Shipping/estimate ?zip ?weight) ?shipCost]
         [(<= ?price ?shipCost)]]
```

get product facts  
needed *during query*

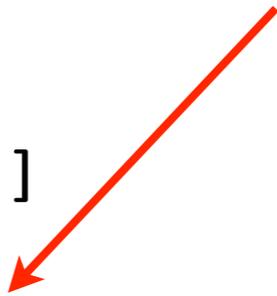


# calling a function

“Find me the customer/product combinations where the shipping cost dominates the product cost.”

```
[ :find ?customer ?product
  :where [?customer :shipAddress ?addr]
         [?addr :zip ?zip]
         [?product :product/weight ?weight]
         [?product :product/price ?price]
         [(Shipping/estimate ?zip ?weight) ?shipCost]
         [(<= ?price ?shipCost)]]
```

call web service  
to bind shipCost



# BYO functions

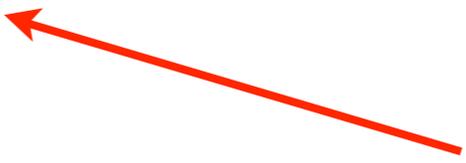
*Functions can be plain  
JVM code.*

```
public class Shipping {  
    public static BigDecimal  
    estimate(String zip1, int pounds);  
}
```

# calling a function

“Find me the customer/product combinations where the shipping cost dominates the product cost.”

```
[ :find ?customer ?product
  :where [?customer :shipAddress ?addr]
         [?addr :zip ?zip]
         [?product :product/weight ?weight]
         [?product :product/price ?price]
         [(Shipping/estimate ?zip ?weight) ?shipCost]
         [(<= ?price ?shipCost)]]
```



constrain price

# calling a function

“Find me the customer/product combinations where the shipping cost dominates the product cost.”

```
[ :find ?customer ?product  
  :where [?customer :shipAddress ?addr]  
         [?addr :zip ?zip]  
         [?product :product/weight ?weight]  
         [?product :product/price ?price]  
         [(Shipping/estimate ?zip ?weight) ?shipCost]  
         [(<= ?price ?shipCost)]]
```

← return customer,  
product pairs

# calling a function

“Find me the customer/product combinations where the shipping cost dominates the product cost.”

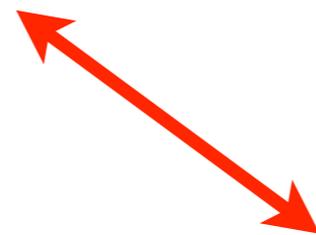
```
[ :find ?customer ?product
  :where [?customer :shipAddress ?addr]
         [?addr :zip ?zip]
         [?product :product/weight ?weight]
         [?product :product/price ?price]
         [(Shipping/estimate ?zip ?weight) ?shipCost]
         [(<= ?price ?shipCost)]]
```

# entities

maplike, point-in-time view of datoms sharing a common **e**

```
{:db/id 42  
 :likes "pizza"  
 :firstName "John"  
 :lastName "Doe"}
```

entity



datoms

```
[42 :likes "pizza"]  
[42 :firstName "John"]  
[42 :lastName "Doe"]
```

# entities

transformation is purely mechanical

```
{ :db/id 42 ← special key for e  
  :likes "pizza"  
  :firstName "John"  
  :lastName "Doe" }
```

```
[ 42 :likes "pizza" ]  
[ 42 :firstName "John" ]  
[ 42 :lastName "Doe" ]
```

# one database, many indexes

<b>structure</b>	<b>attribute</b>
k/v	AVET
row	EAVT
column	AEVT
document	EAVT, partitions, components
graph	VAET

# lab: query

explore mbrainz data with Query API	<a href="https://github.com/Datomic/mbrainz-sample/wiki/Queries">https://github.com/Datomic/mbrainz-sample/wiki/Queries</a>
explore mbrainz data with Pull API	<a href="http://docs.datomic.com/pull.html">http://docs.datomic.com/pull.html</a>
query your own data	<a href="http://docs.datomic.com/query.html">http://docs.datomic.com/query.html</a>
pull your own data	<a href="http://docs.datomic.com/pull.html">http://docs.datomic.com/pull.html</a>
navigate your data with the Entity API	<a href="http://docs.datomic.com/entities.html">http://docs.datomic.com/entities.html</a>

# time model

reified transactions

t & basis-t

log

filters

sync

excision

# reified transactions

entities like any other entity in the system

associated with every datom

increasing entity ids over time

associated with increasing counter **t**

have a `:db/txInstant`

have any other attributes you specify

have their own index (the log)

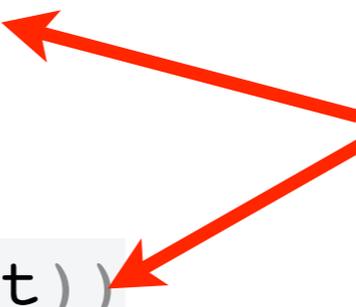
# basis

```
(def basis-t (d/basis-t db))  
=> 130223
```

```
(def basis-tx (d/t->tx basis-t))  
=> 13194139663535
```

```
(d/pull db '[*] basis-tx)  
=> {:db/id 13194139663535,  
    :db/txInstant #inst "2014-10-13T22:58:03.832-00:00"}
```

basis is most recent t, tx



facts about basis tx

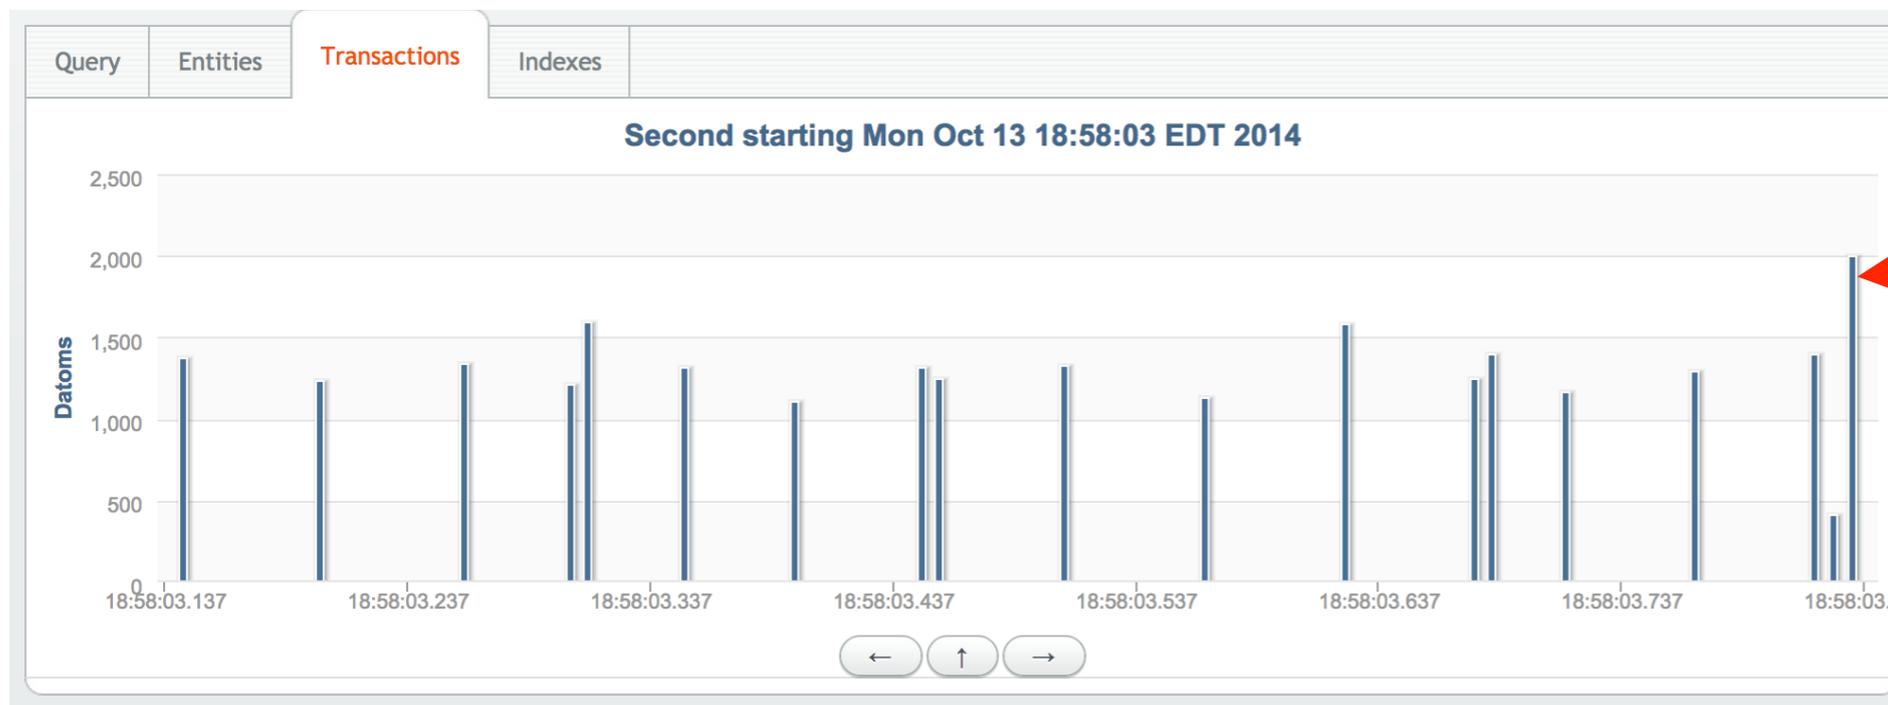


a *time point* is any of: t, tx, instant

# log

```
(def log (d/log conn))  
(-> (d/tx-range log basis-tx (inc basis-tx))  
    seq first :data count)
```

contents of txes



basis tx

# transaction attributes

tx partition  
references current tx

```
[{:db/id #db/id [:db.part/user]  
 :story/title "codeq"  
 :story/url "http://blog.datomic.com/2012/10/codeq.html"}  
{:db/id (d/tempid :db.part/tx)  
 :publish/at (java.util.Date.)}]
```

add your own  
attributes

```
[{:db/id [:item/id "DLC-042"]  
 :item/count 250}  
{:db/id #db/id [:db.part/tx]  
 :db/txInstant #inst "2013-02"}]
```

override txInstant  
(for imports)

<https://github.com/Datomic/day-of-datomic/blob/master/tutorial/filter.clj>

<https://github.com/Datomic/day-of-datomic/blob/master/tutorial/filters.clj>

# filters

filter	semantics	supports
<i>(default)</i>	current state	what is the current situation?
<code>.asOf</code>	state at point in past	how were things in the past?
<code>.since</code>	state since point in past	how have things changed?
<code>.history</code>	timeless view of all history	anything!

# raw data

<b>e</b>	<b>a</b>	<b>v</b>	<b>tx</b>	<b>added</b>
0x0000c00000003e9	:db/txInstant	Mon Dec 31 19:00:00	0xc00000003e9	TRUE
0x00001000000003ea	:item/id	DLC-042	0xc00000003e9	TRUE
0x00001000000003ea	:item/description	Dilitihium Crystals	0xc00000003e9	TRUE
0x00001000000003ea	:item/count	100	0xc00000003e9	TRUE
0x0000c00000003eb	:db/txInstant	Thu Jan 31 19:00:00	0xc00000003eb	TRUE
0x00001000000003ea	:item/count	100	0xc00000003eb	FALSE
0x00001000000003ea	:item/count	250	0xc00000003eb	TRUE
0x0000c00000003ec	:db/txInstant	Thu Feb 27 19:00:00	0xc00000003ec	TRUE
0x00001000000003ea	:item/count	250	0xc00000003ec	FALSE
0x00001000000003ea	:item/count	50	0xc00000003ec	TRUE
0x0000c00000003ed	:db/txInstant	Mon Mar 31 20:00:00	0xc00000003ed	TRUE
0x0000c00000003ed	:tx/error	TRUE	0xc00000003ed	TRUE
0x00001000000003ea	:item/count	50	0xc00000003ed	FALSE
0x00001000000003ea	:item/count	9999	0xc00000003ed	TRUE
0x0000c00000003ee	:db/txInstant	Wed May 14 20:00:00	0xc00000003ee	TRUE
0x00001000000003ea	:item/count	9999	0xc00000003ee	FALSE
0x00001000000003ea	:item/count	100	0xc00000003ee	TRUE

# default filter

<b>e</b>	<b>a</b>	<b>v</b>	<b>tx</b>	<b>added</b>
0x0000c00000003e9	:db/txInstant	Mon Dec 31 19:00:00	0xc00000003e9	TRUE
0x00001000000003ea	:item/id	DLC-042	0xc00000003e9	TRUE
0x00001000000003ea	:item/description	Dilitihium Crystals	0xc00000003e9	TRUE
0x00001000000003ea	:item/count	100	0xc00000003e9	TRUE
0x0000c00000003eb	:db/txInstant	Thu Jan 31 19:00:00	0xc00000003eb	TRUE
0x00001000000003ea	:item/count	100	0xc00000003eb	FALSE
0x00001000000003ea	:item/count	250	0xc00000003eb	TRUE
0x0000c00000003ec	:db/txInstant	Thu Feb 27 19:00:00	0xc00000003ec	TRUE
0x00001000000003ea	:item/count	250	0xc00000003ec	FALSE
0x00001000000003ea	:item/count	50	0xc00000003ec	TRUE
0x0000c00000003ed	:db/txInstant	Mon Mar 31 20:00:00	0xc00000003ed	TRUE
0x0000c00000003ed	:tx/error	TRUE	0xc00000003ed	TRUE
0x00001000000003ea	:item/count	50	0xc00000003ed	FALSE
0x00001000000003ea	:item/count	9999	0xc00000003ed	TRUE
0x0000c00000003ee	:db/txInstant	Wed May 14 20:00:00	0xc00000003ee	TRUE
0x00001000000003ea	:item/count	9999	0xc00000003ee	FALSE
0x00001000000003ea	:item/count	100	0xc00000003ee	TRUE

# as-of filter

<b>e</b>	<b>a</b>	<b>v</b>	<b>tx</b>	<b>added</b>
0x0000c00000003e9	:db/txInstant	Mon Dec 31 19:00:00	0xc00000003e9	TRUE
0x00001000000003ea	:item/id	DLC-042	0xc00000003e9	TRUE
0x00001000000003ea	:item/description	Dilitihium Crystals	0xc00000003e9	TRUE
0x00001000000003ea	:item/count	100	0xc00000003e9	TRUE
0x0000c00000003eb	:db/txInstant	Thu Jan 31 19:00:00	0xc00000003eb	TRUE
0x00001000000003ea	:item/count	100	0xc00000003eb	FALSE
0x00001000000003ea	:item/count	250	0xc00000003eb	TRUE
0x0000c00000003ec	:db/txInstant	Thu Feb 27 19:00:00	0xc00000003ec	TRUE
0x00001000000003ea	:item/count	250	0xc00000003ec	FALSE
0x00001000000003ea	:item/count	50	0xc00000003ec	TRUE
0x0000c00000003ed	:db/txInstant	Mon Mar 31 20:00:00	0xc00000003ed	TRUE
0x0000c00000003ed	:tx/error	TRUE	0xc00000003ed	TRUE
0x00001000000003ea	:item/count	50	0xc00000003ed	FALSE
0x00001000000003ea	:item/count	9999	0xc00000003ed	TRUE
0x0000c00000003ee	:db/txInstant	Wed May 14 20:00:00	0xc00000003ee	TRUE
0x00001000000003ea	:item/count	9999	0xc00000003ee	FALSE
0x00001000000003ea	:item/count	100	0xc00000003ee	TRUE

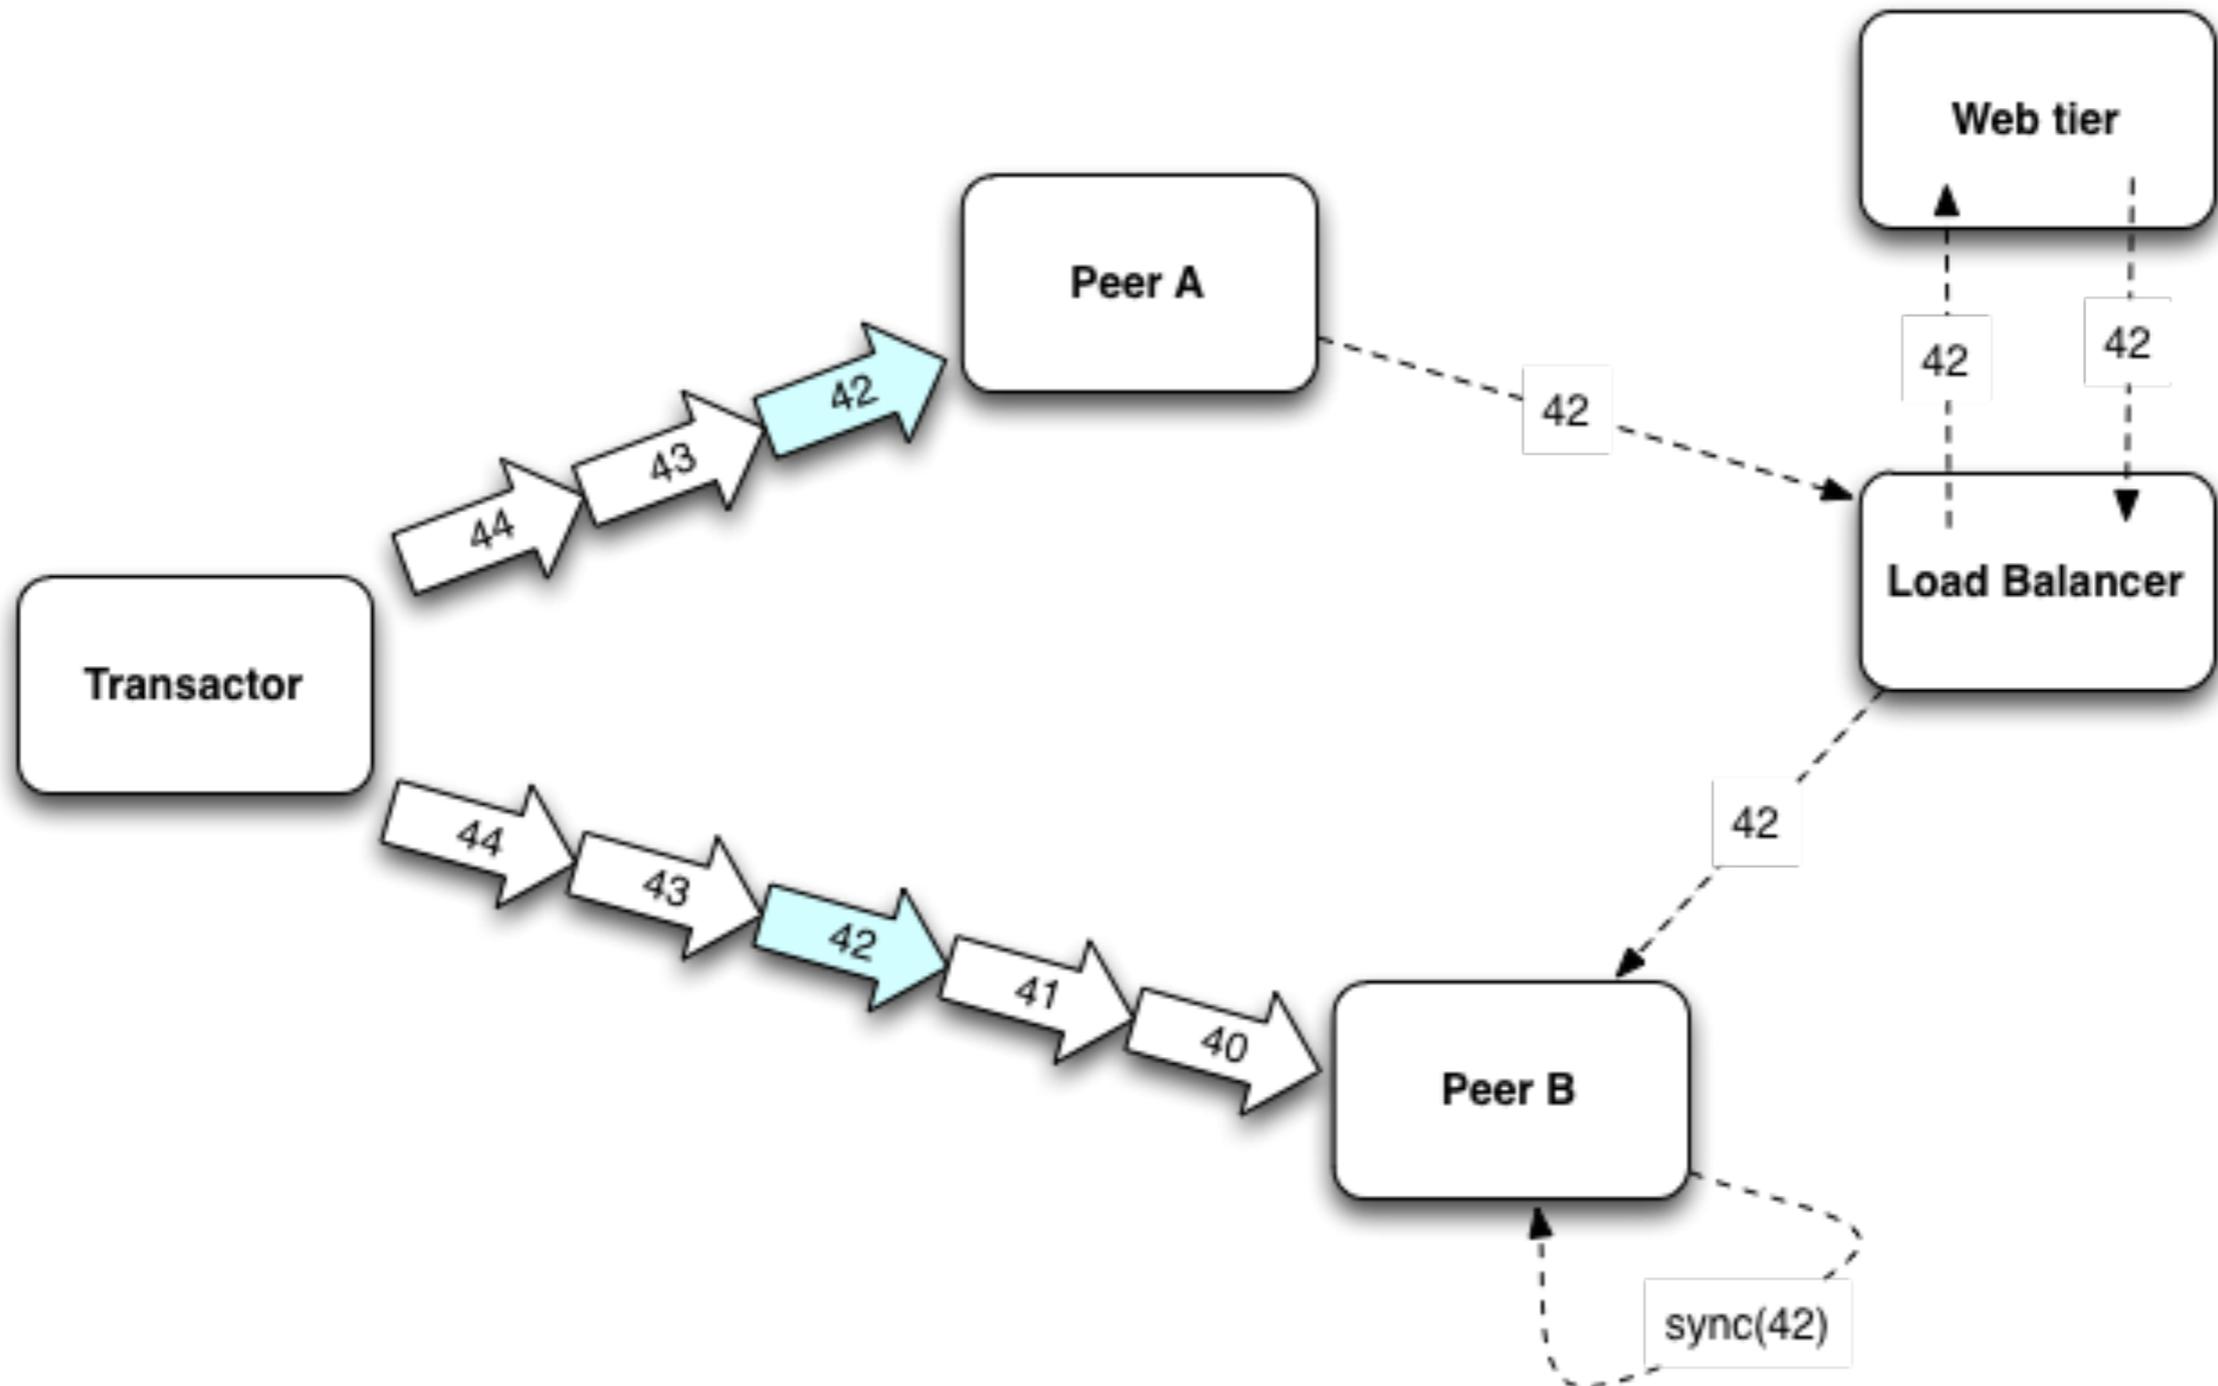
# since filter

<b>e</b>	<b>a</b>	<b>v</b>	<b>tx</b>	<b>added</b>
0x0000c00000003e9	:db/txInstant	Mon Dec 31 19:00:00	0xc00000003e9	TRUE
0x00001000000003ea	:item/id	DLC-042	0xc00000003e9	TRUE
0x00001000000003ea	:item/description	Dilitihium Crystals	0xc00000003e9	TRUE
0x00001000000003ea	:item/count	100	0xc00000003e9	TRUE
0x0000c00000003eb	:db/txInstant	Thu Jan 31 19:00:00	0xc00000003eb	TRUE
0x00001000000003ea	:item/count	100	0xc00000003eb	FALSE
0x00001000000003ea	:item/count	250	0xc00000003eb	TRUE
0x0000c00000003ec	:db/txInstant	Thu Feb 27 19:00:00	0xc00000003ec	TRUE
0x00001000000003ea	:item/count	250	0xc00000003ec	FALSE
0x00001000000003ea	:item/count	50	0xc00000003ec	TRUE
0x0000c00000003ed	:db/txInstant	Mon Mar 31 20:00:00	0xc00000003ed	TRUE
0x0000c00000003ed	:tx/error	TRUE	0xc00000003ed	TRUE
0x00001000000003ea	:item/count	50	0xc00000003ed	FALSE
0x00001000000003ea	:item/count	9999	0xc00000003ed	TRUE
0x0000c00000003ee	:db/txInstant	Wed May 14 20:00:00	0xc00000003ee	TRUE
0x00001000000003ea	:item/count	9999	0xc00000003ee	FALSE
0x00001000000003ea	:item/count	100	0xc00000003ee	TRUE

# history = raw

<b>e</b>	<b>a</b>	<b>v</b>	<b>tx</b>	<b>added</b>
0x0000c00000003e9	:db/txInstant	Mon Dec 31 19:00:00	0xc00000003e9	TRUE
0x00001000000003ea	:item/id	DLC-042	0xc00000003e9	TRUE
0x00001000000003ea	:item/description	Dilitihium Crystals	0xc00000003e9	TRUE
0x00001000000003ea	:item/count	100	0xc00000003e9	TRUE
0x0000c00000003eb	:db/txInstant	Thu Jan 31 19:00:00	0xc00000003eb	TRUE
0x00001000000003ea	:item/count	100	0xc00000003eb	FALSE
0x00001000000003ea	:item/count	250	0xc00000003eb	TRUE
0x0000c00000003ec	:db/txInstant	Thu Feb 27 19:00:00	0xc00000003ec	TRUE
0x00001000000003ea	:item/count	250	0xc00000003ec	FALSE
0x00001000000003ea	:item/count	50	0xc00000003ec	TRUE
0x0000c00000003ed	:db/txInstant	Mon Mar 31 20:00:00	0xc00000003ed	TRUE
0x0000c00000003ed	:tx/error	TRUE	0xc00000003ed	TRUE
0x00001000000003ea	:item/count	50	0xc00000003ed	FALSE
0x00001000000003ea	:item/count	9999	0xc00000003ed	TRUE
0x0000c00000003ee	:db/txInstant	Wed May 14 20:00:00	0xc00000003ee	TRUE
0x00001000000003ea	:item/count	9999	0xc00000003ee	FALSE
0x00001000000003ea	:item/count	100	0xc00000003ee	TRUE

# sync



excision: permanently,  
unrecoverably, lose data

don't do this  
unless you must

# lab: navigating time

explore your data  
with filters in console

explore your data  
with transaction view in console

query against an as-of db

use the log to recover a recent  
transaction

annotate a transaction

<https://github.com/Datomic/day-of-datomic/blob/master/tutorial/filters.clj>

[https://github.com/Datomic/day-of-datomic/blob/master/tutorial/basis\\_t\\_and\\_log.clj](https://github.com/Datomic/day-of-datomic/blob/master/tutorial/basis_t_and_log.clj)

<https://github.com/Datomic/day-of-datomic/blob/master/tutorial/filters.clj>

# operational model

caching

indexing

bring your own storage

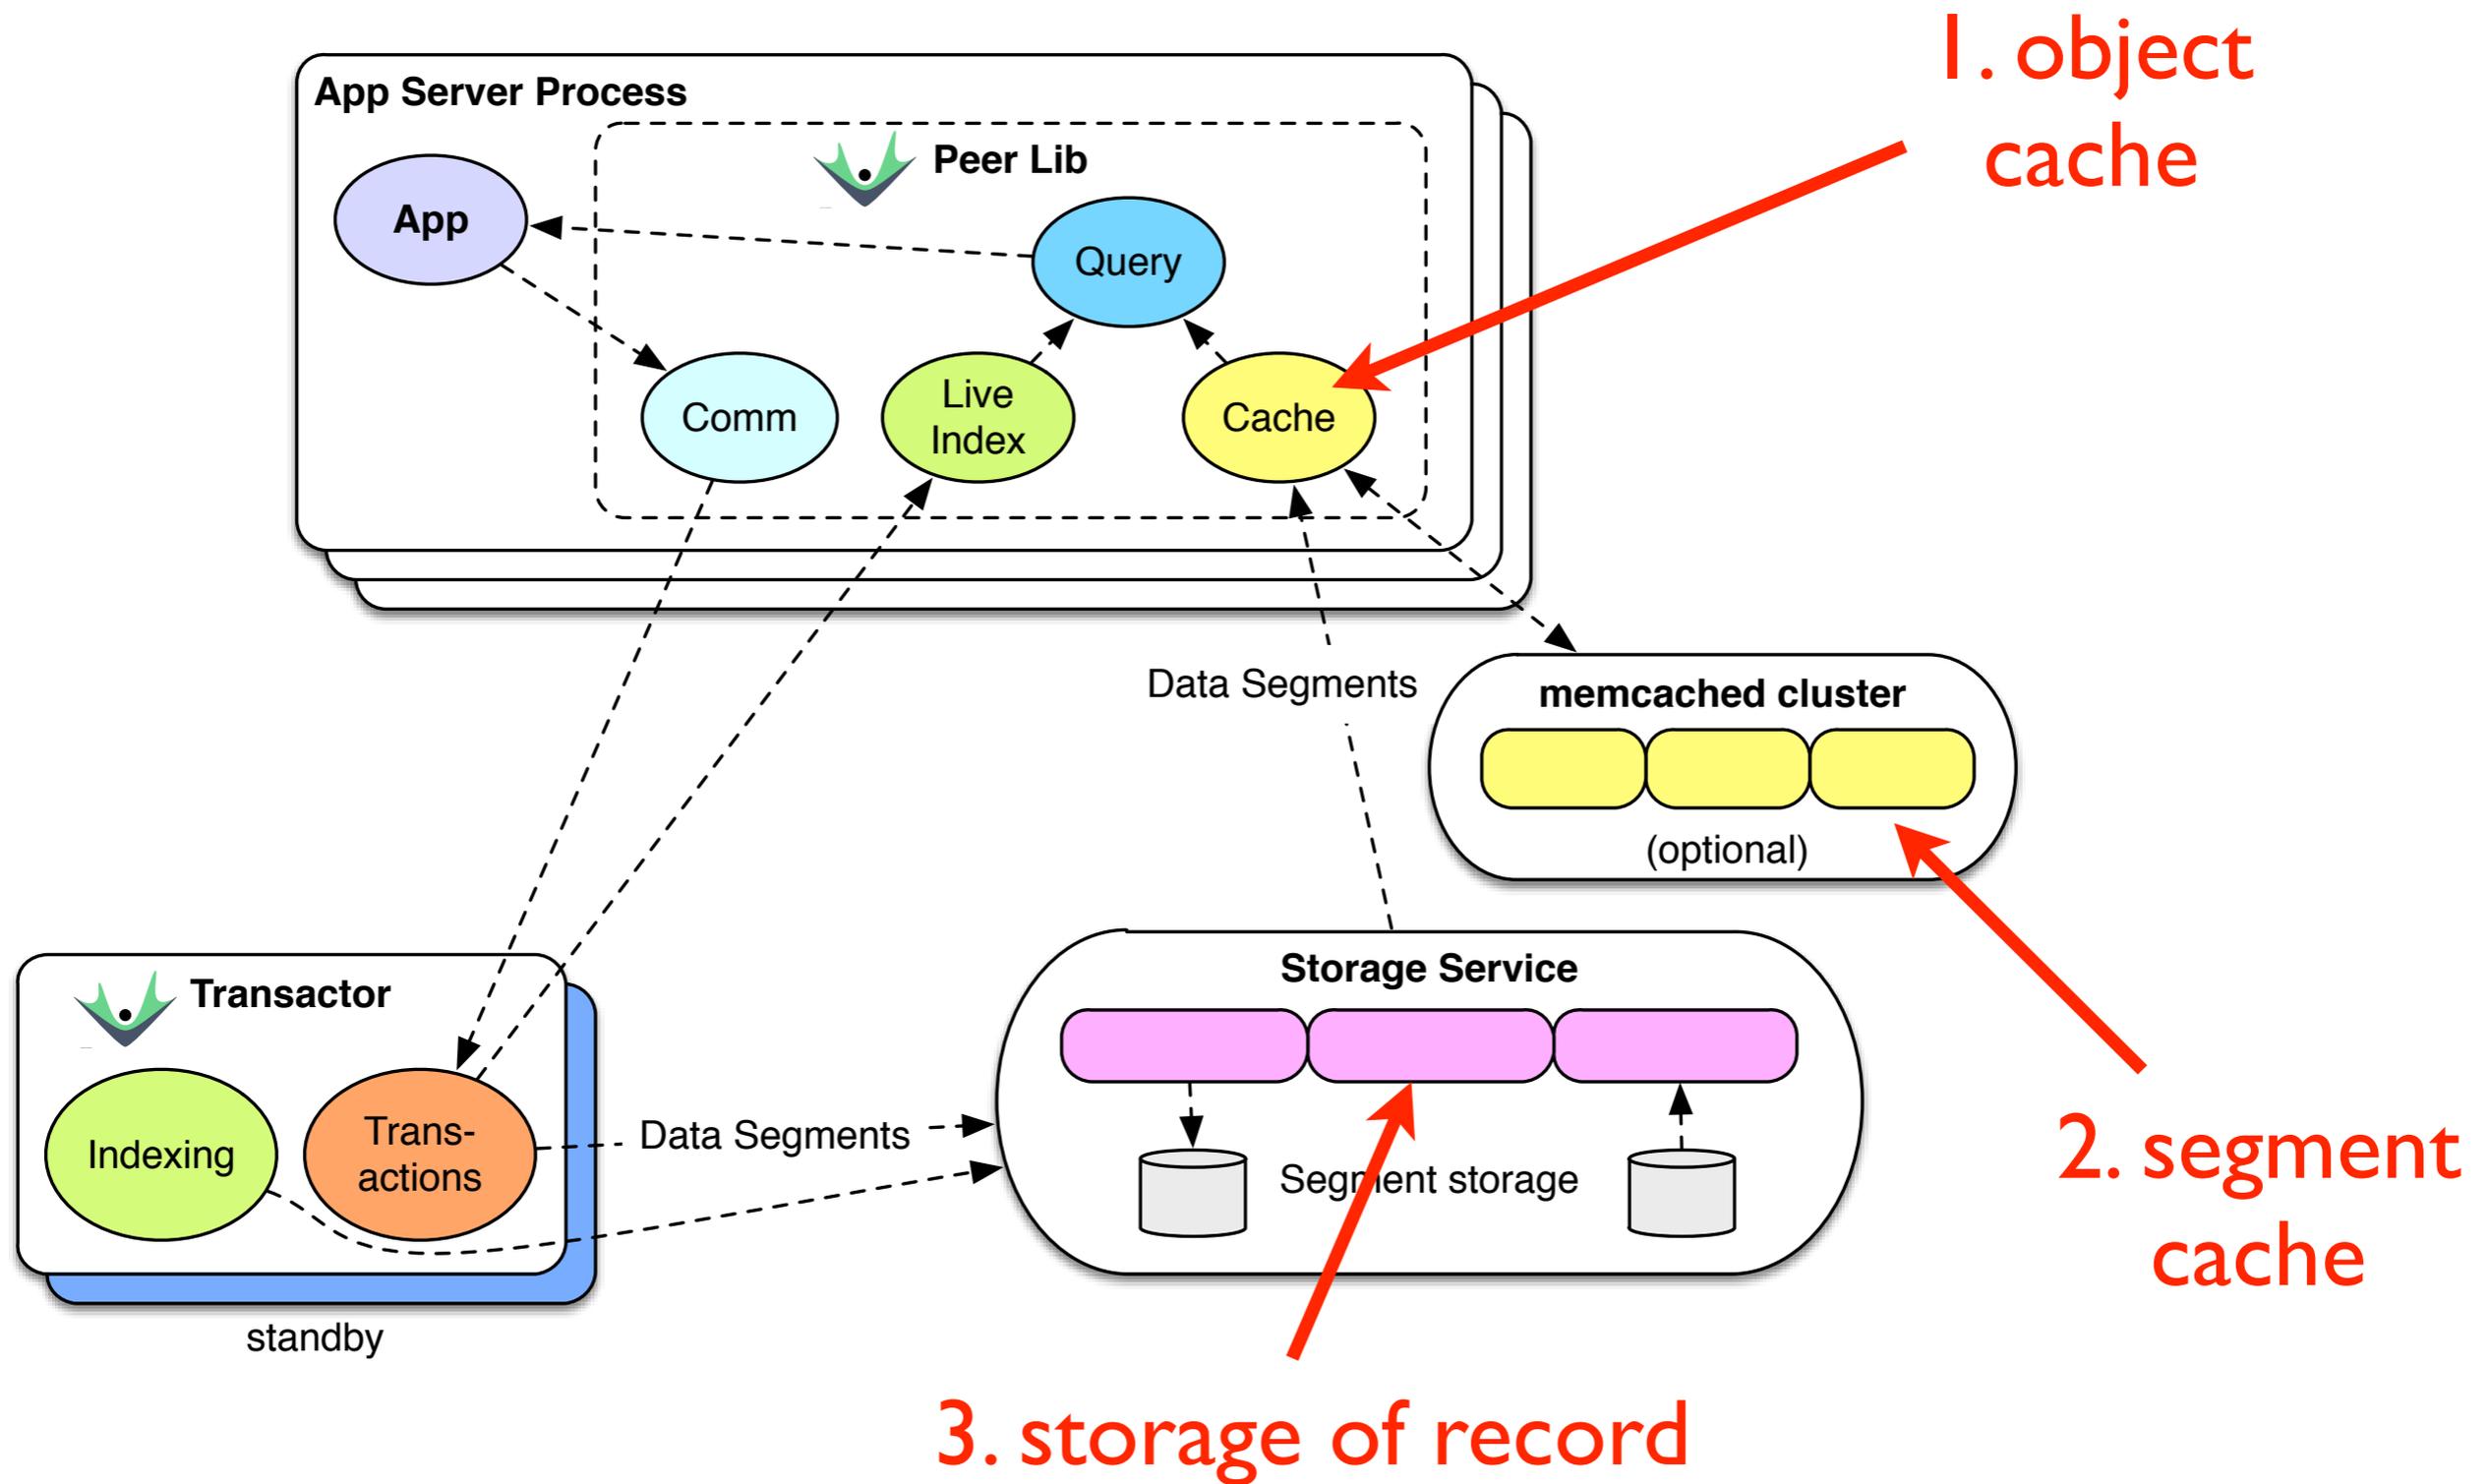
transactors

peers

monitoring

capacity planning

# caching



# understanding caching

single knob: how much?

*segments* (1000s of datoms), not individual facts

segments are immutable, never expire

segment names can cohabit with other caches

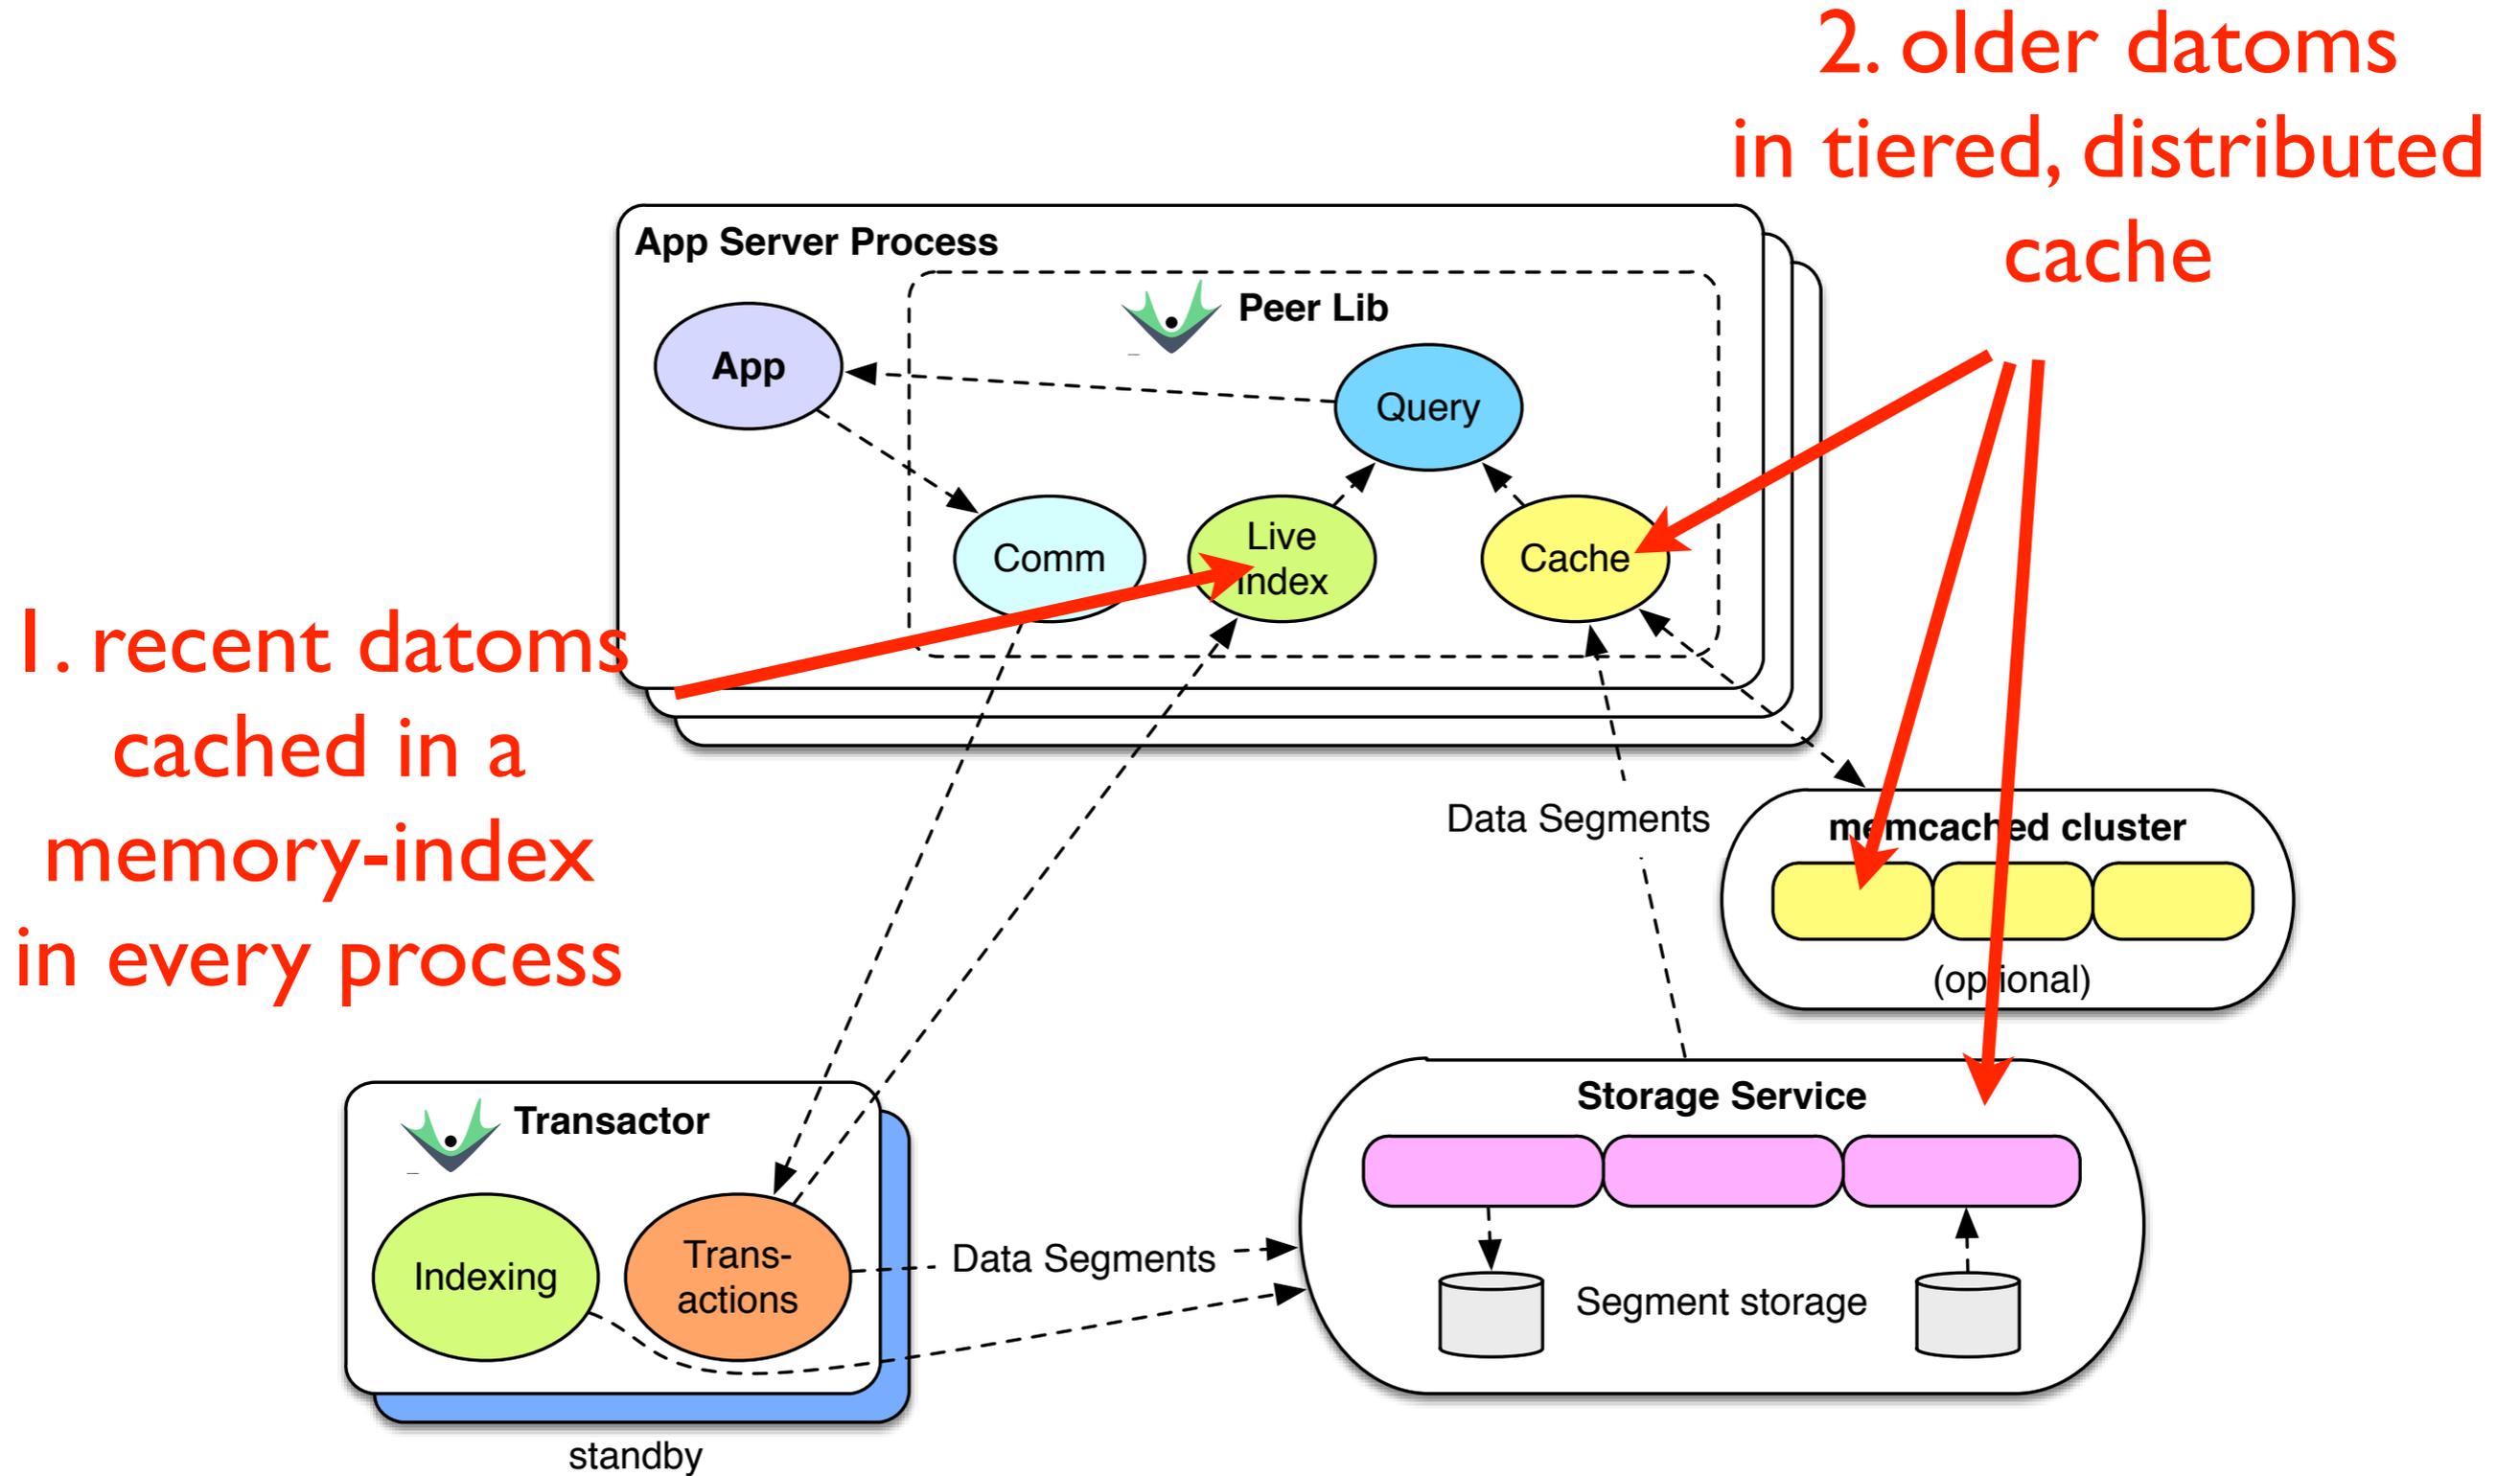
object cache is *in your process!*

will beat any RPC-based database when cache hits

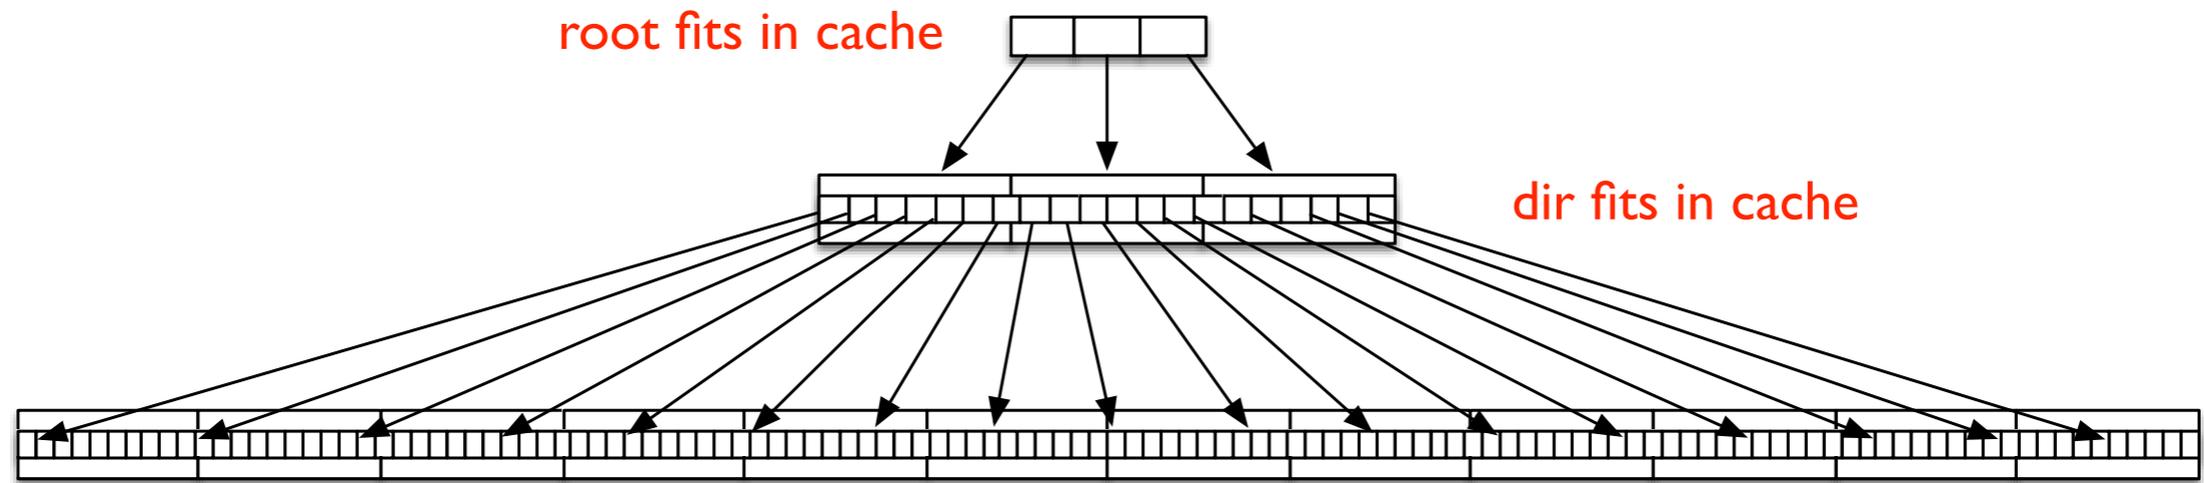
# reading data

source	format	capacity	latency
peer object cache	Java objects	~ 10 GB / peer	sub-microsecond
memcached	zipped fressian	arbitrary	~ 1 msec
storage	zipped fressian	arbitrary	~ 10 msec

# indexing



# index trees



segments fit in cache for small dbs,  
1-2 reads away for larger dbs

*history is in separate trees, and is  
never in the way of present queries*

# understanding indexing

indexes made as needed

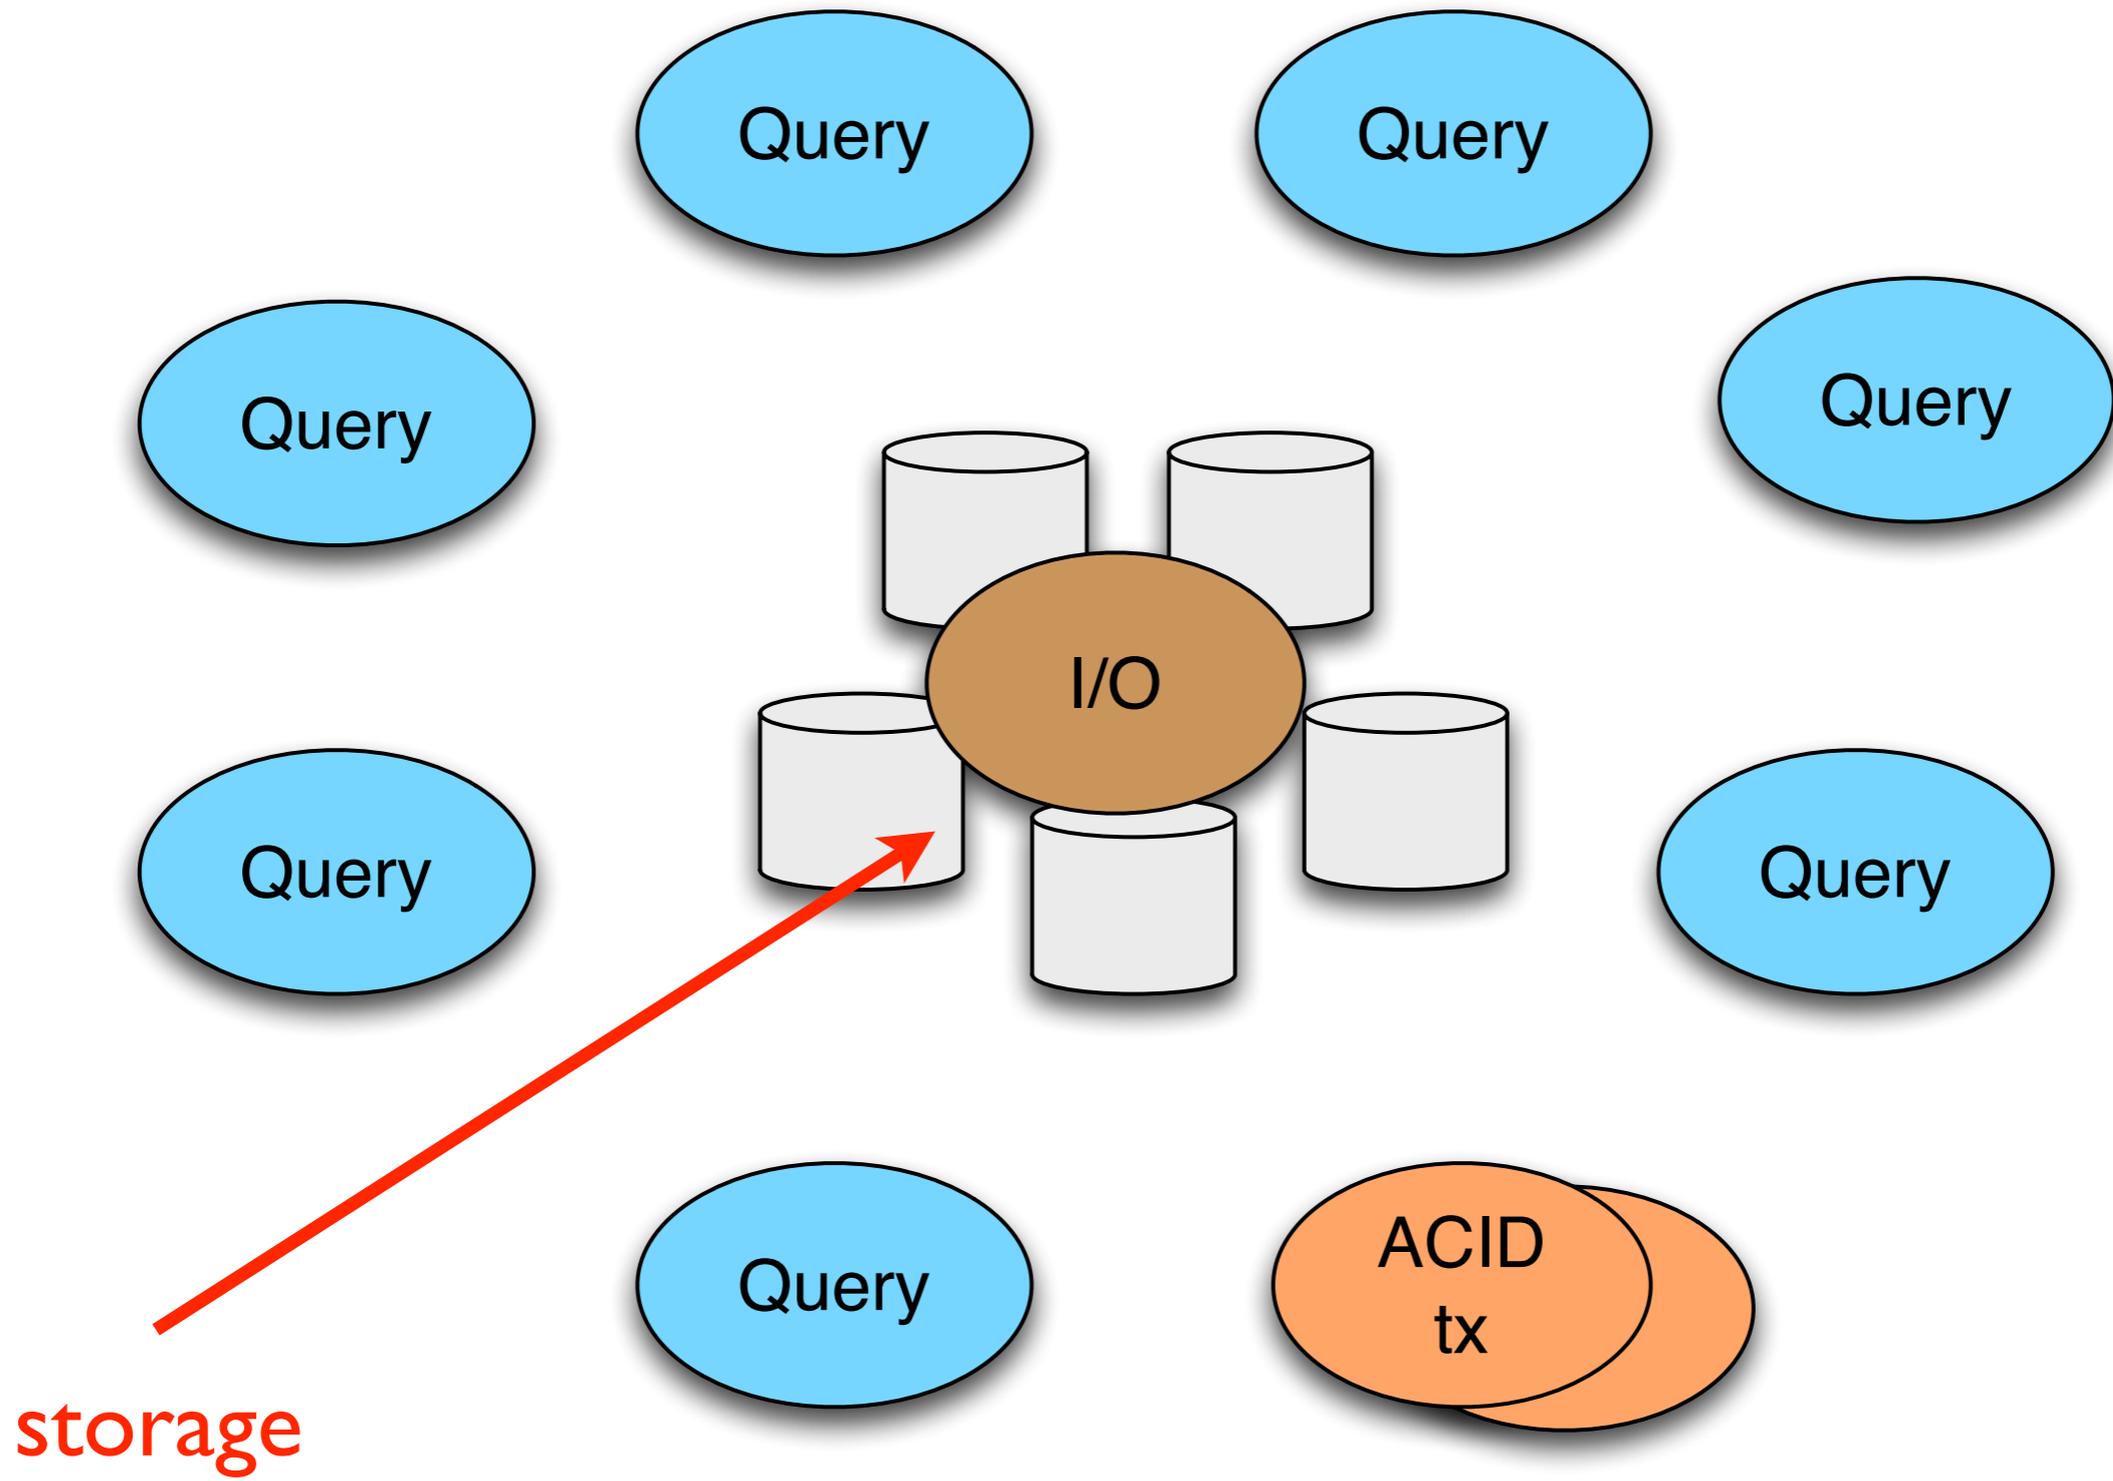
background process (on transactor)

indexing can fall too far behind

memory index exceeds predefined threshold

transactor throttles writes

only likely to see during import jobs



storage

# storage options

dynamodb

sql (any JDBC)

cassandra

riak

couchbase

infinispan

dev (local disk)

# choosing storage

more important

reliability

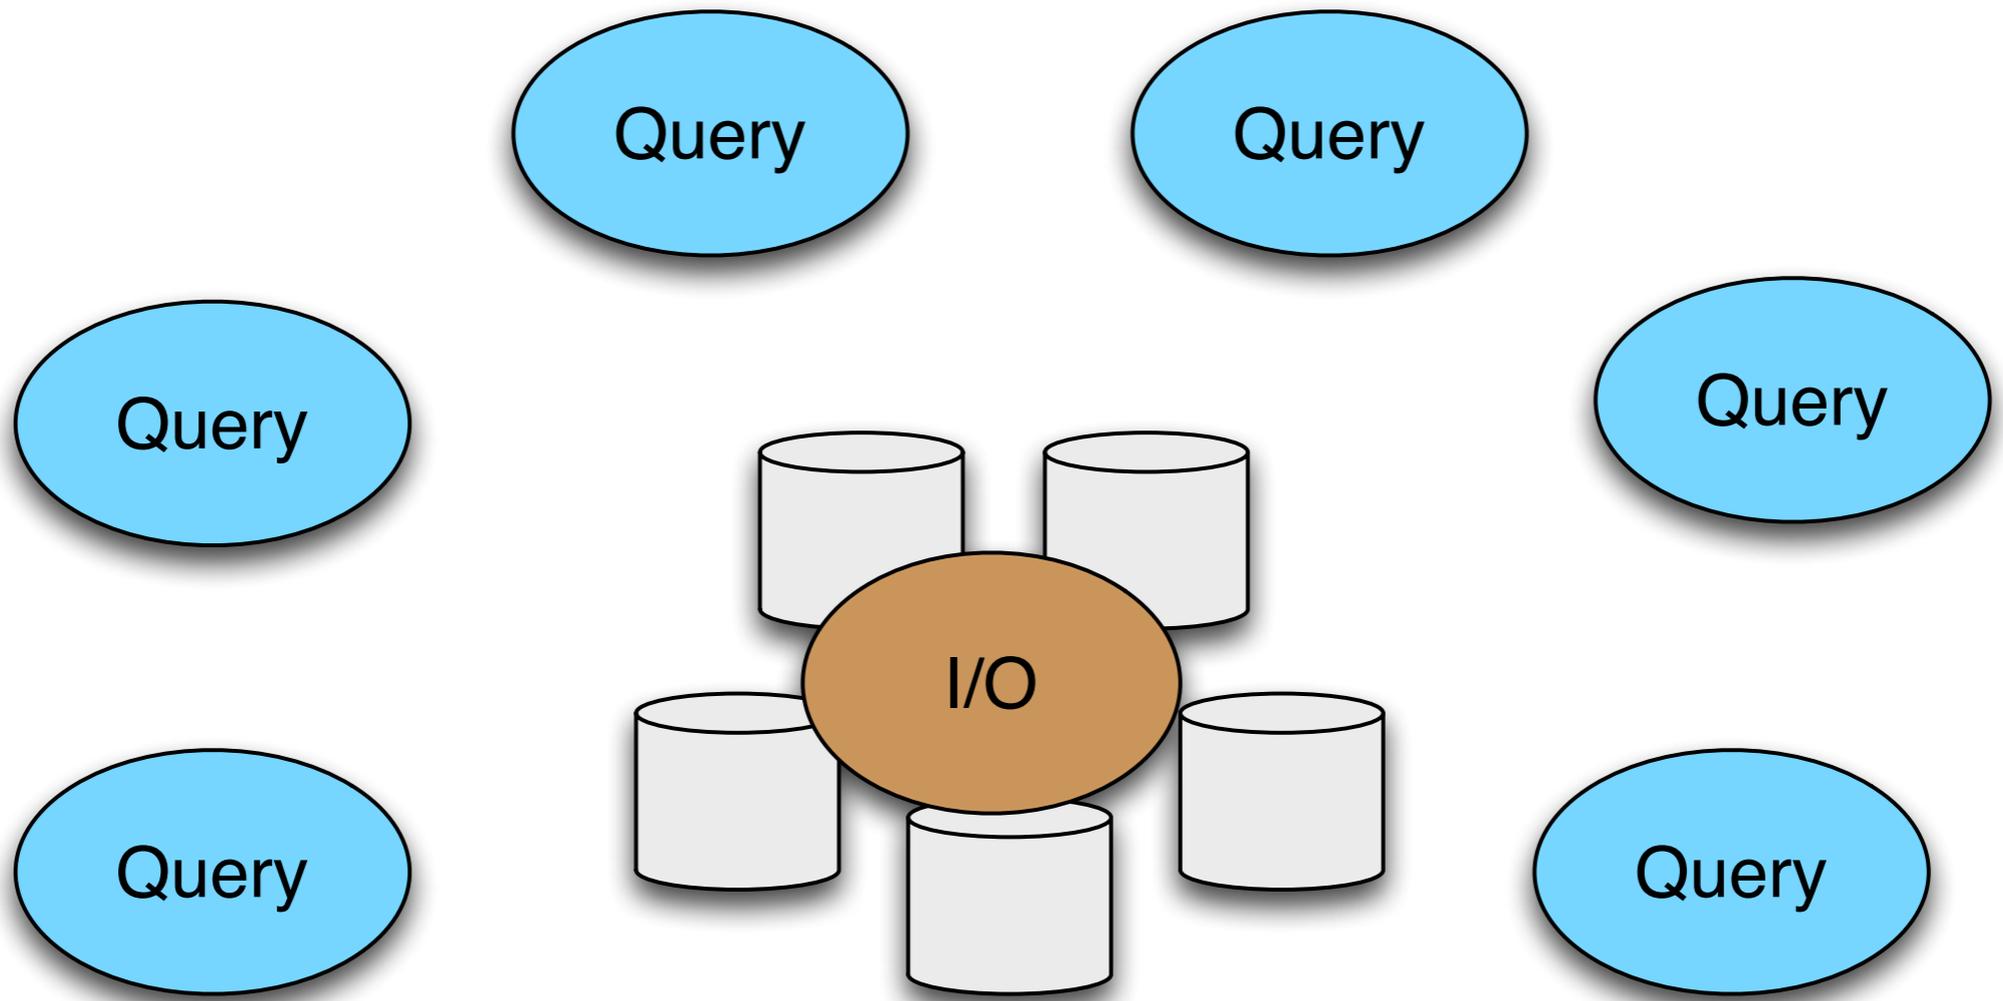
manageability

familiarity

less important

datomic insulates storage from load

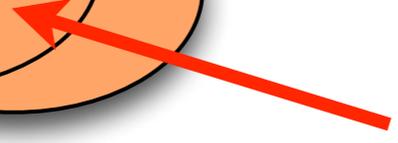
datomic tolerates storage latency



Query

ACID  
tx

transactors



# understanding transactors

ACID (single writer thread!)

stream recent change to peers

manage indexing

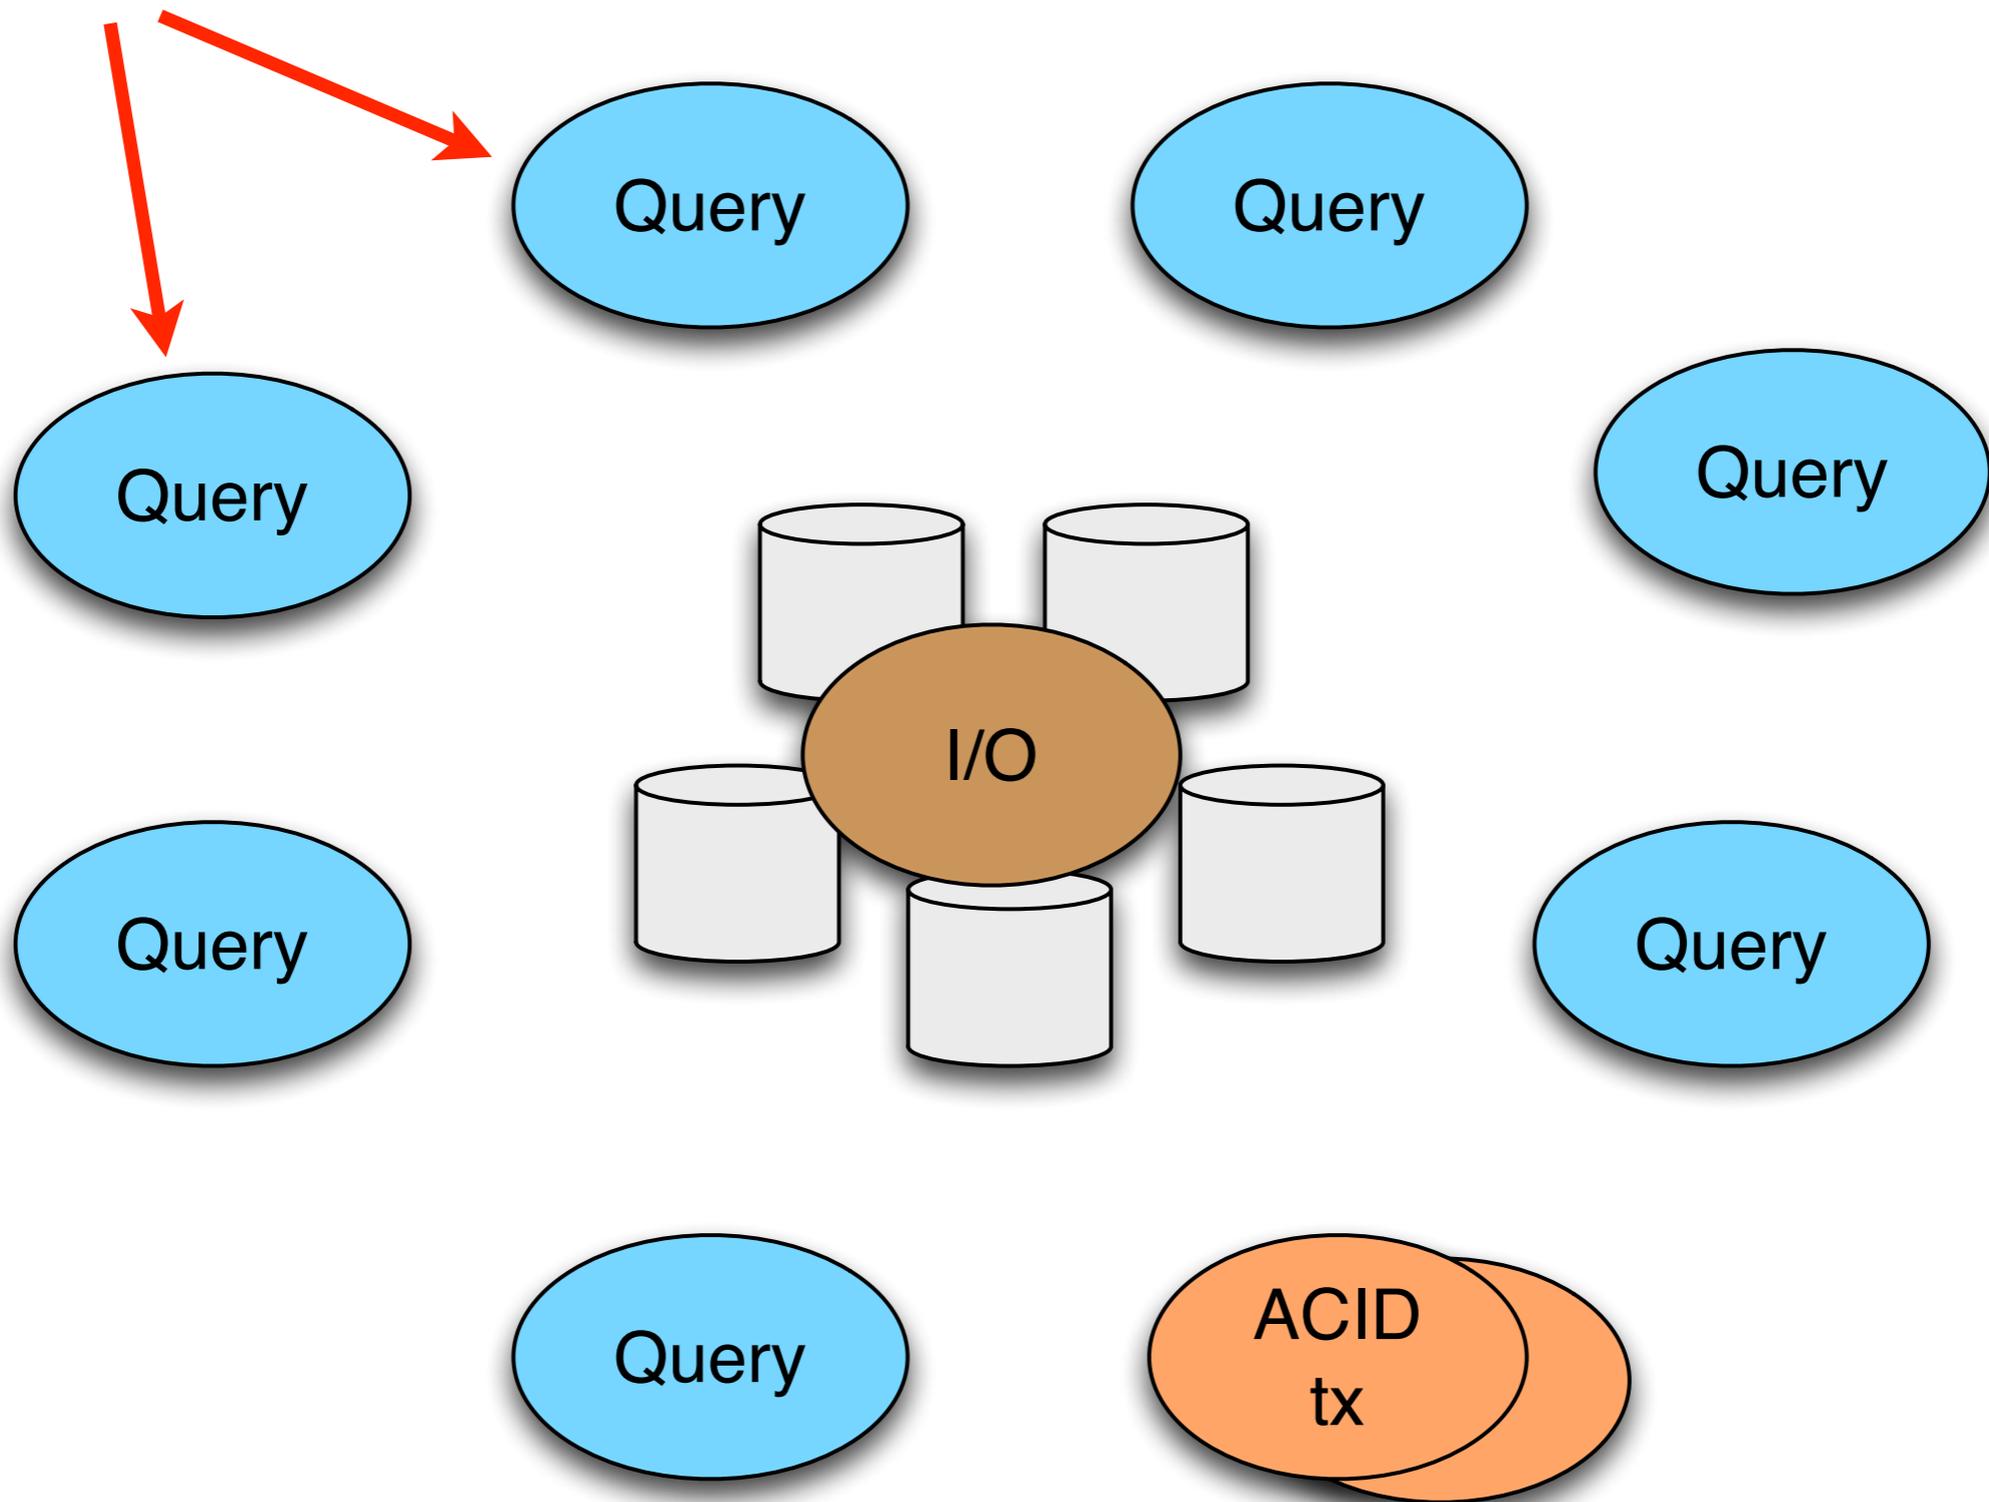
dead-simple HA configuration

point two machines at one storage

dead-simple HA coordination

conditional puts on storage

peers



# understanding peers

embedded JVM lib

directly in application servers (most common)

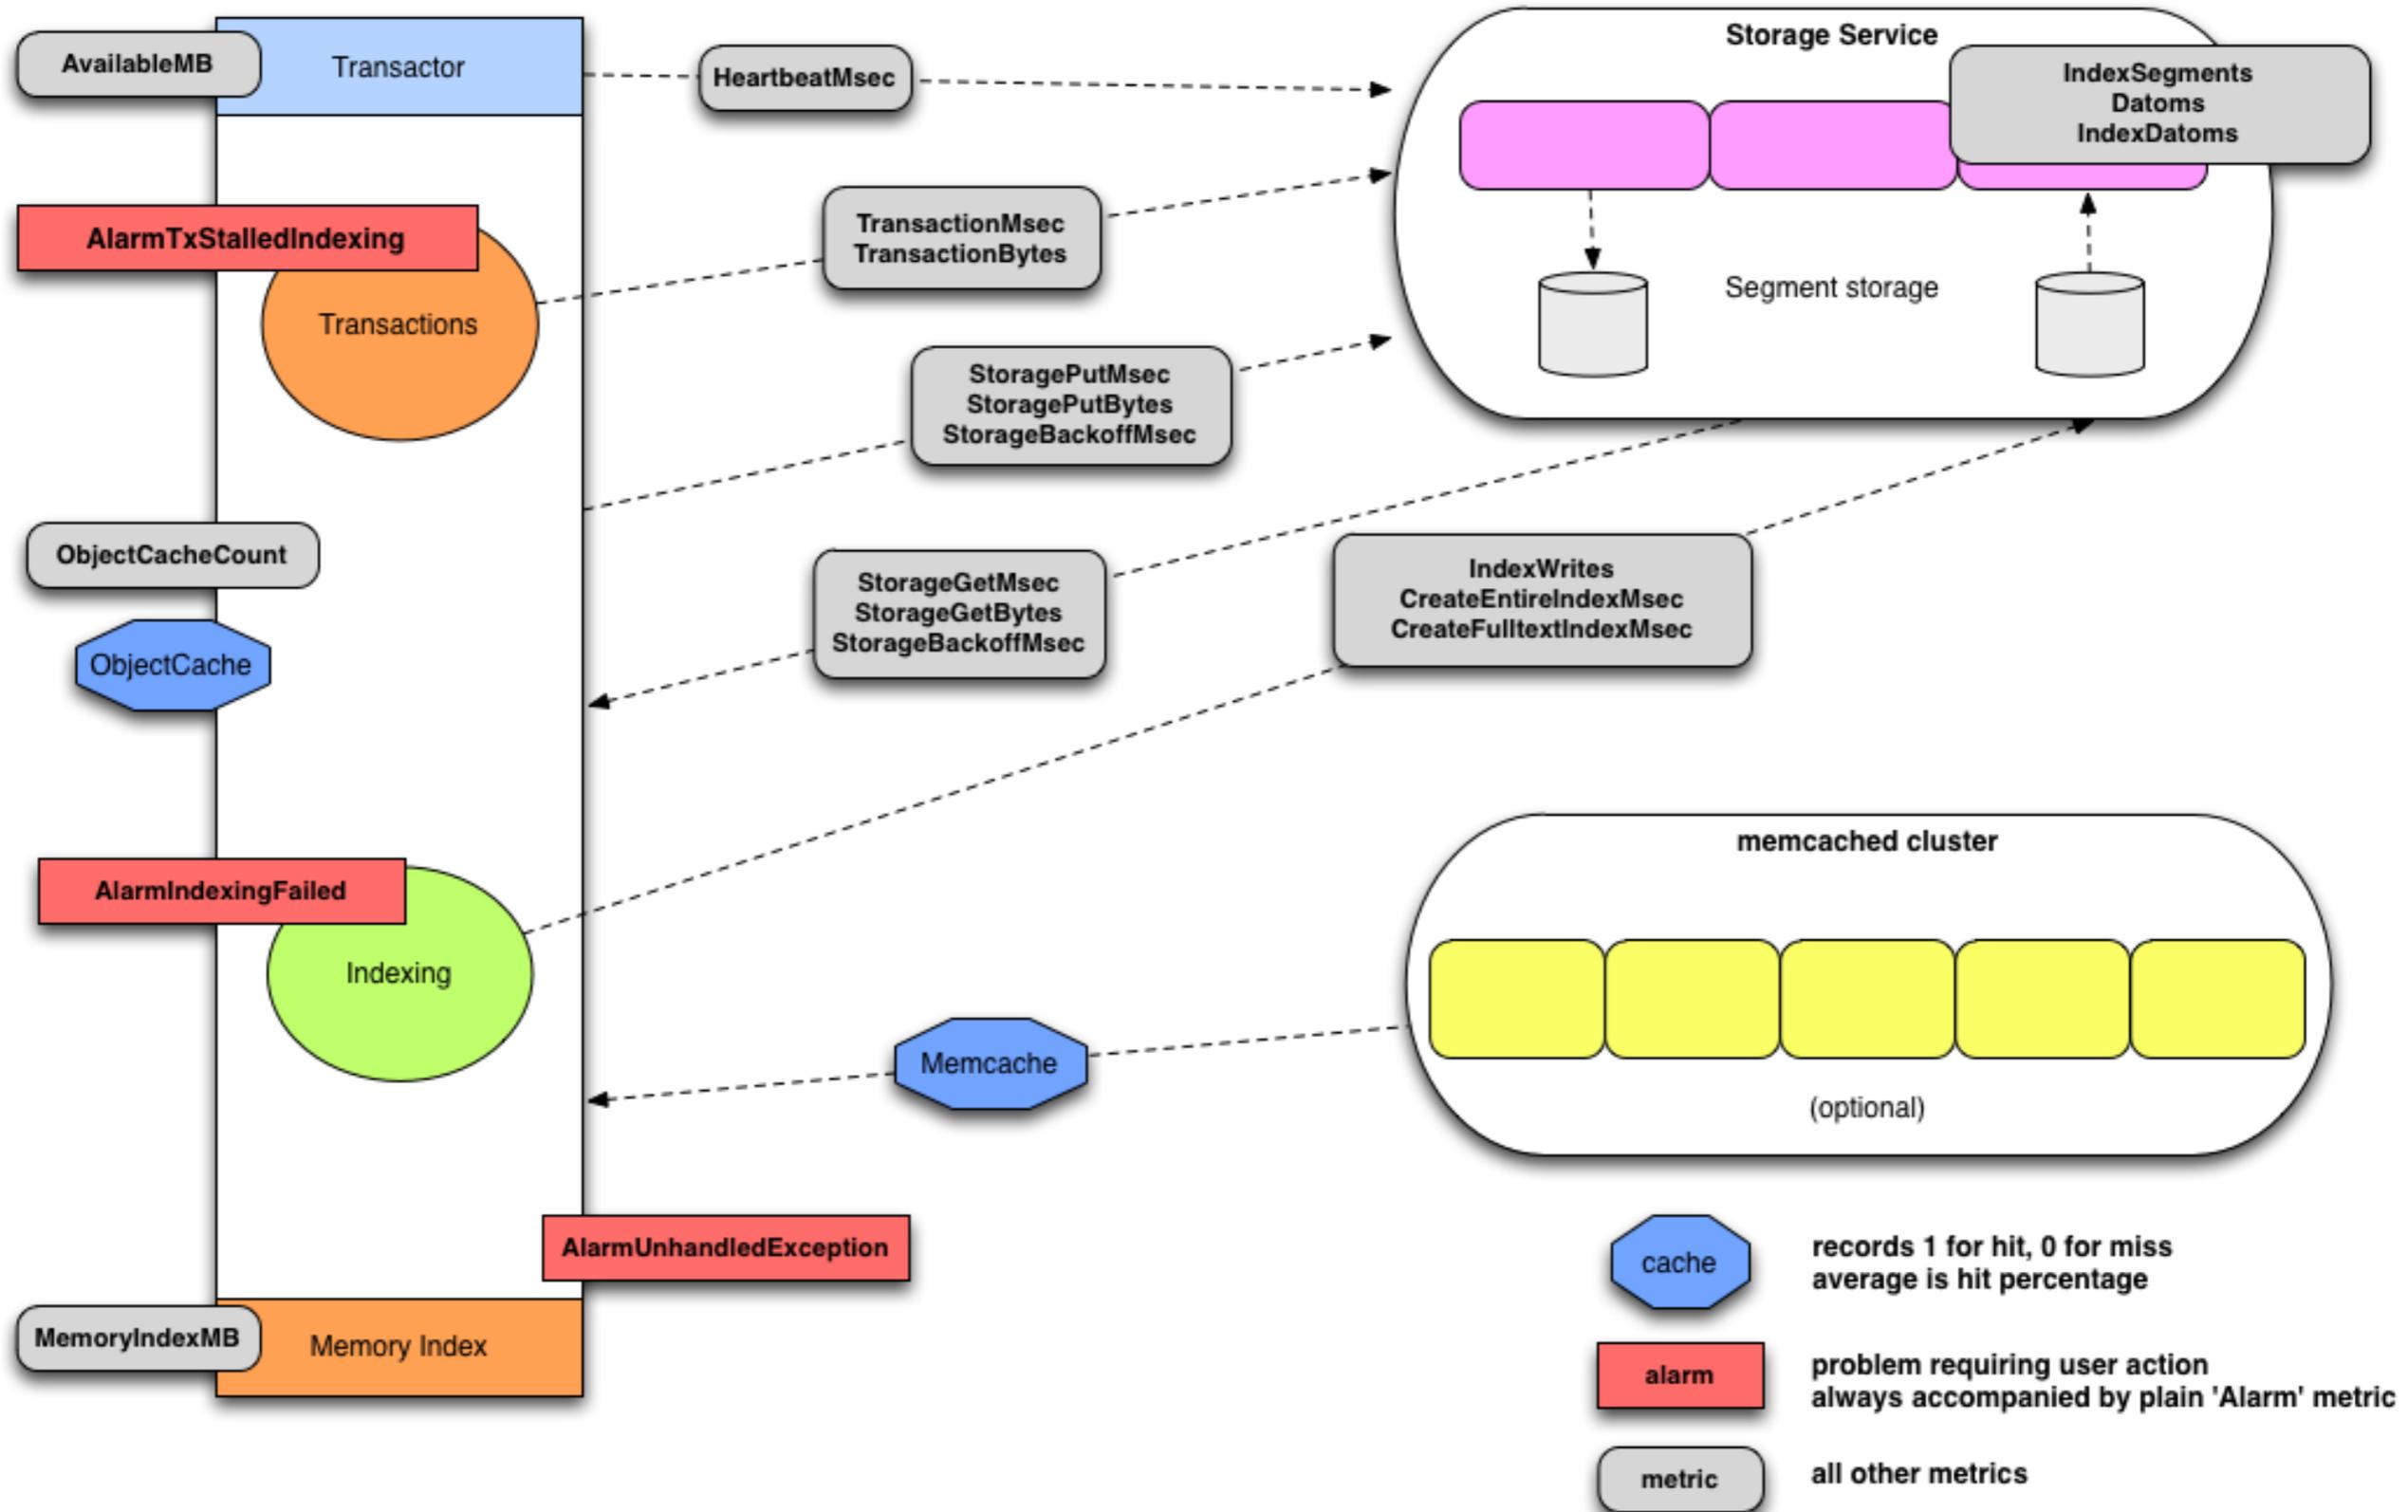
introduce a service tier

query load does not bother the transactor

answer queries in-memory via automatic cache

`datomic.objectCacheMax` (default 50% RAM)

# Datomic Monitoring



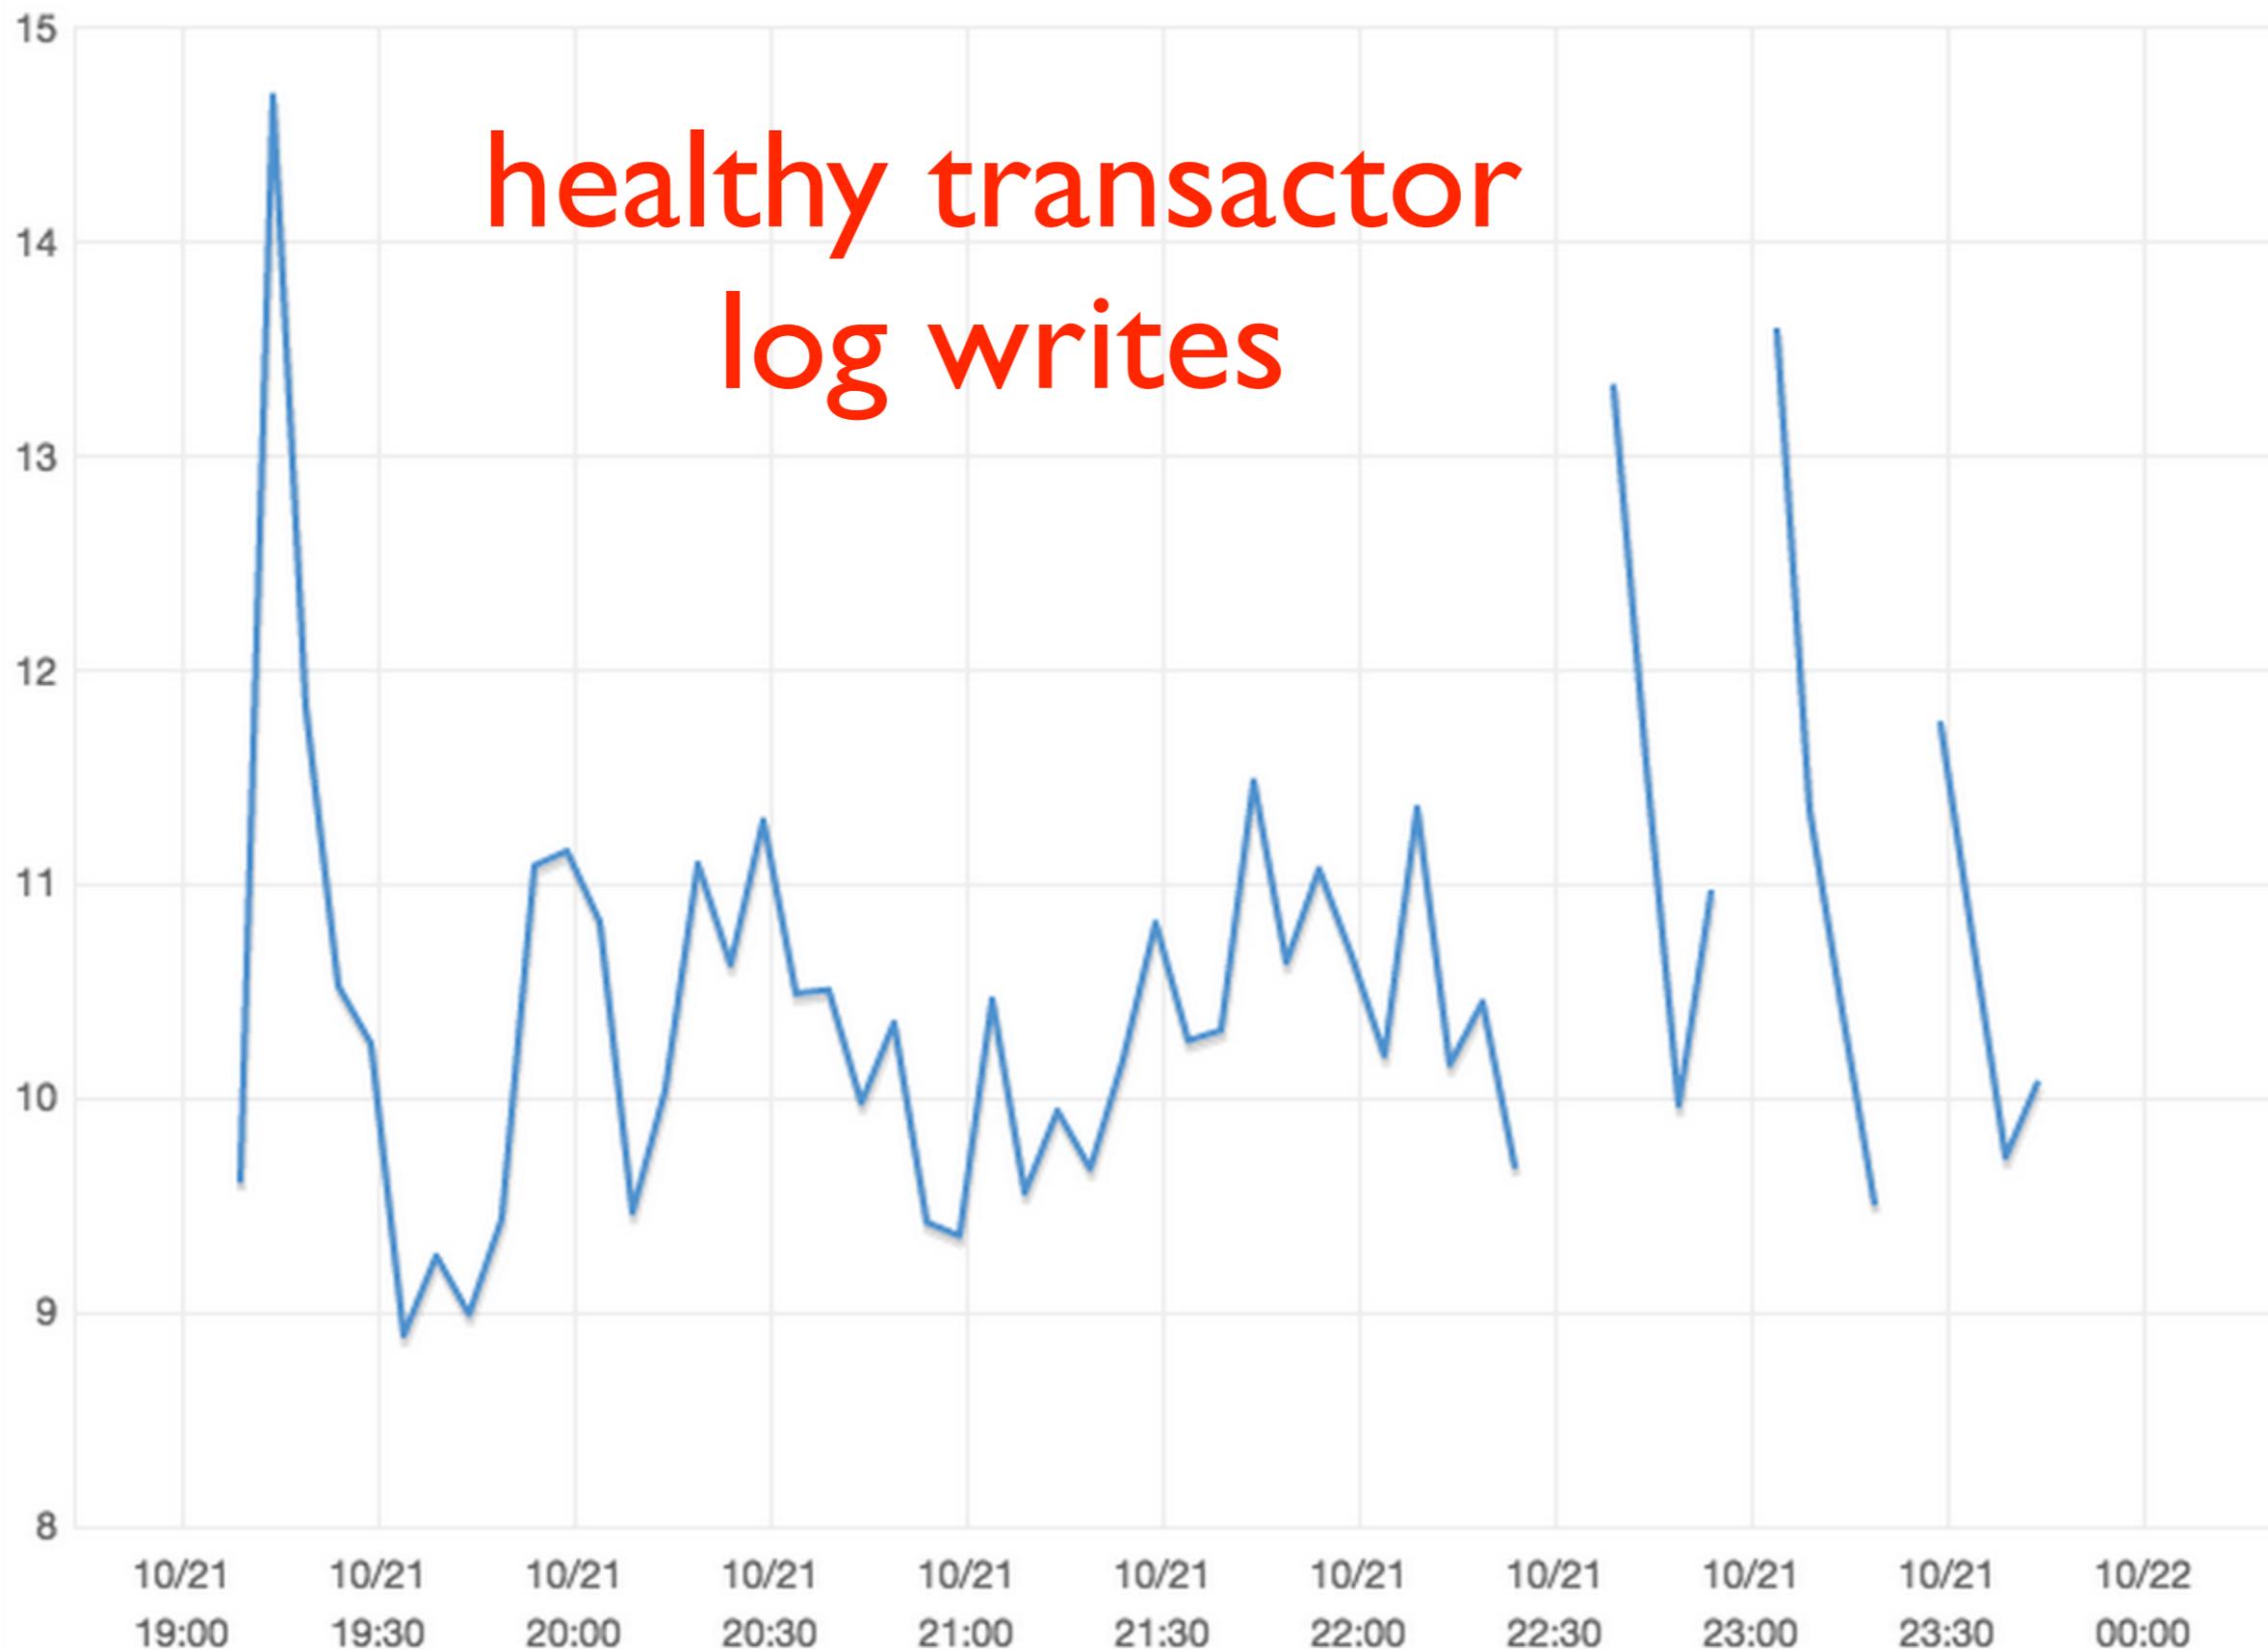
LogWriteMsec (Milliseconds)

Average ▾

5 Minutes ▾



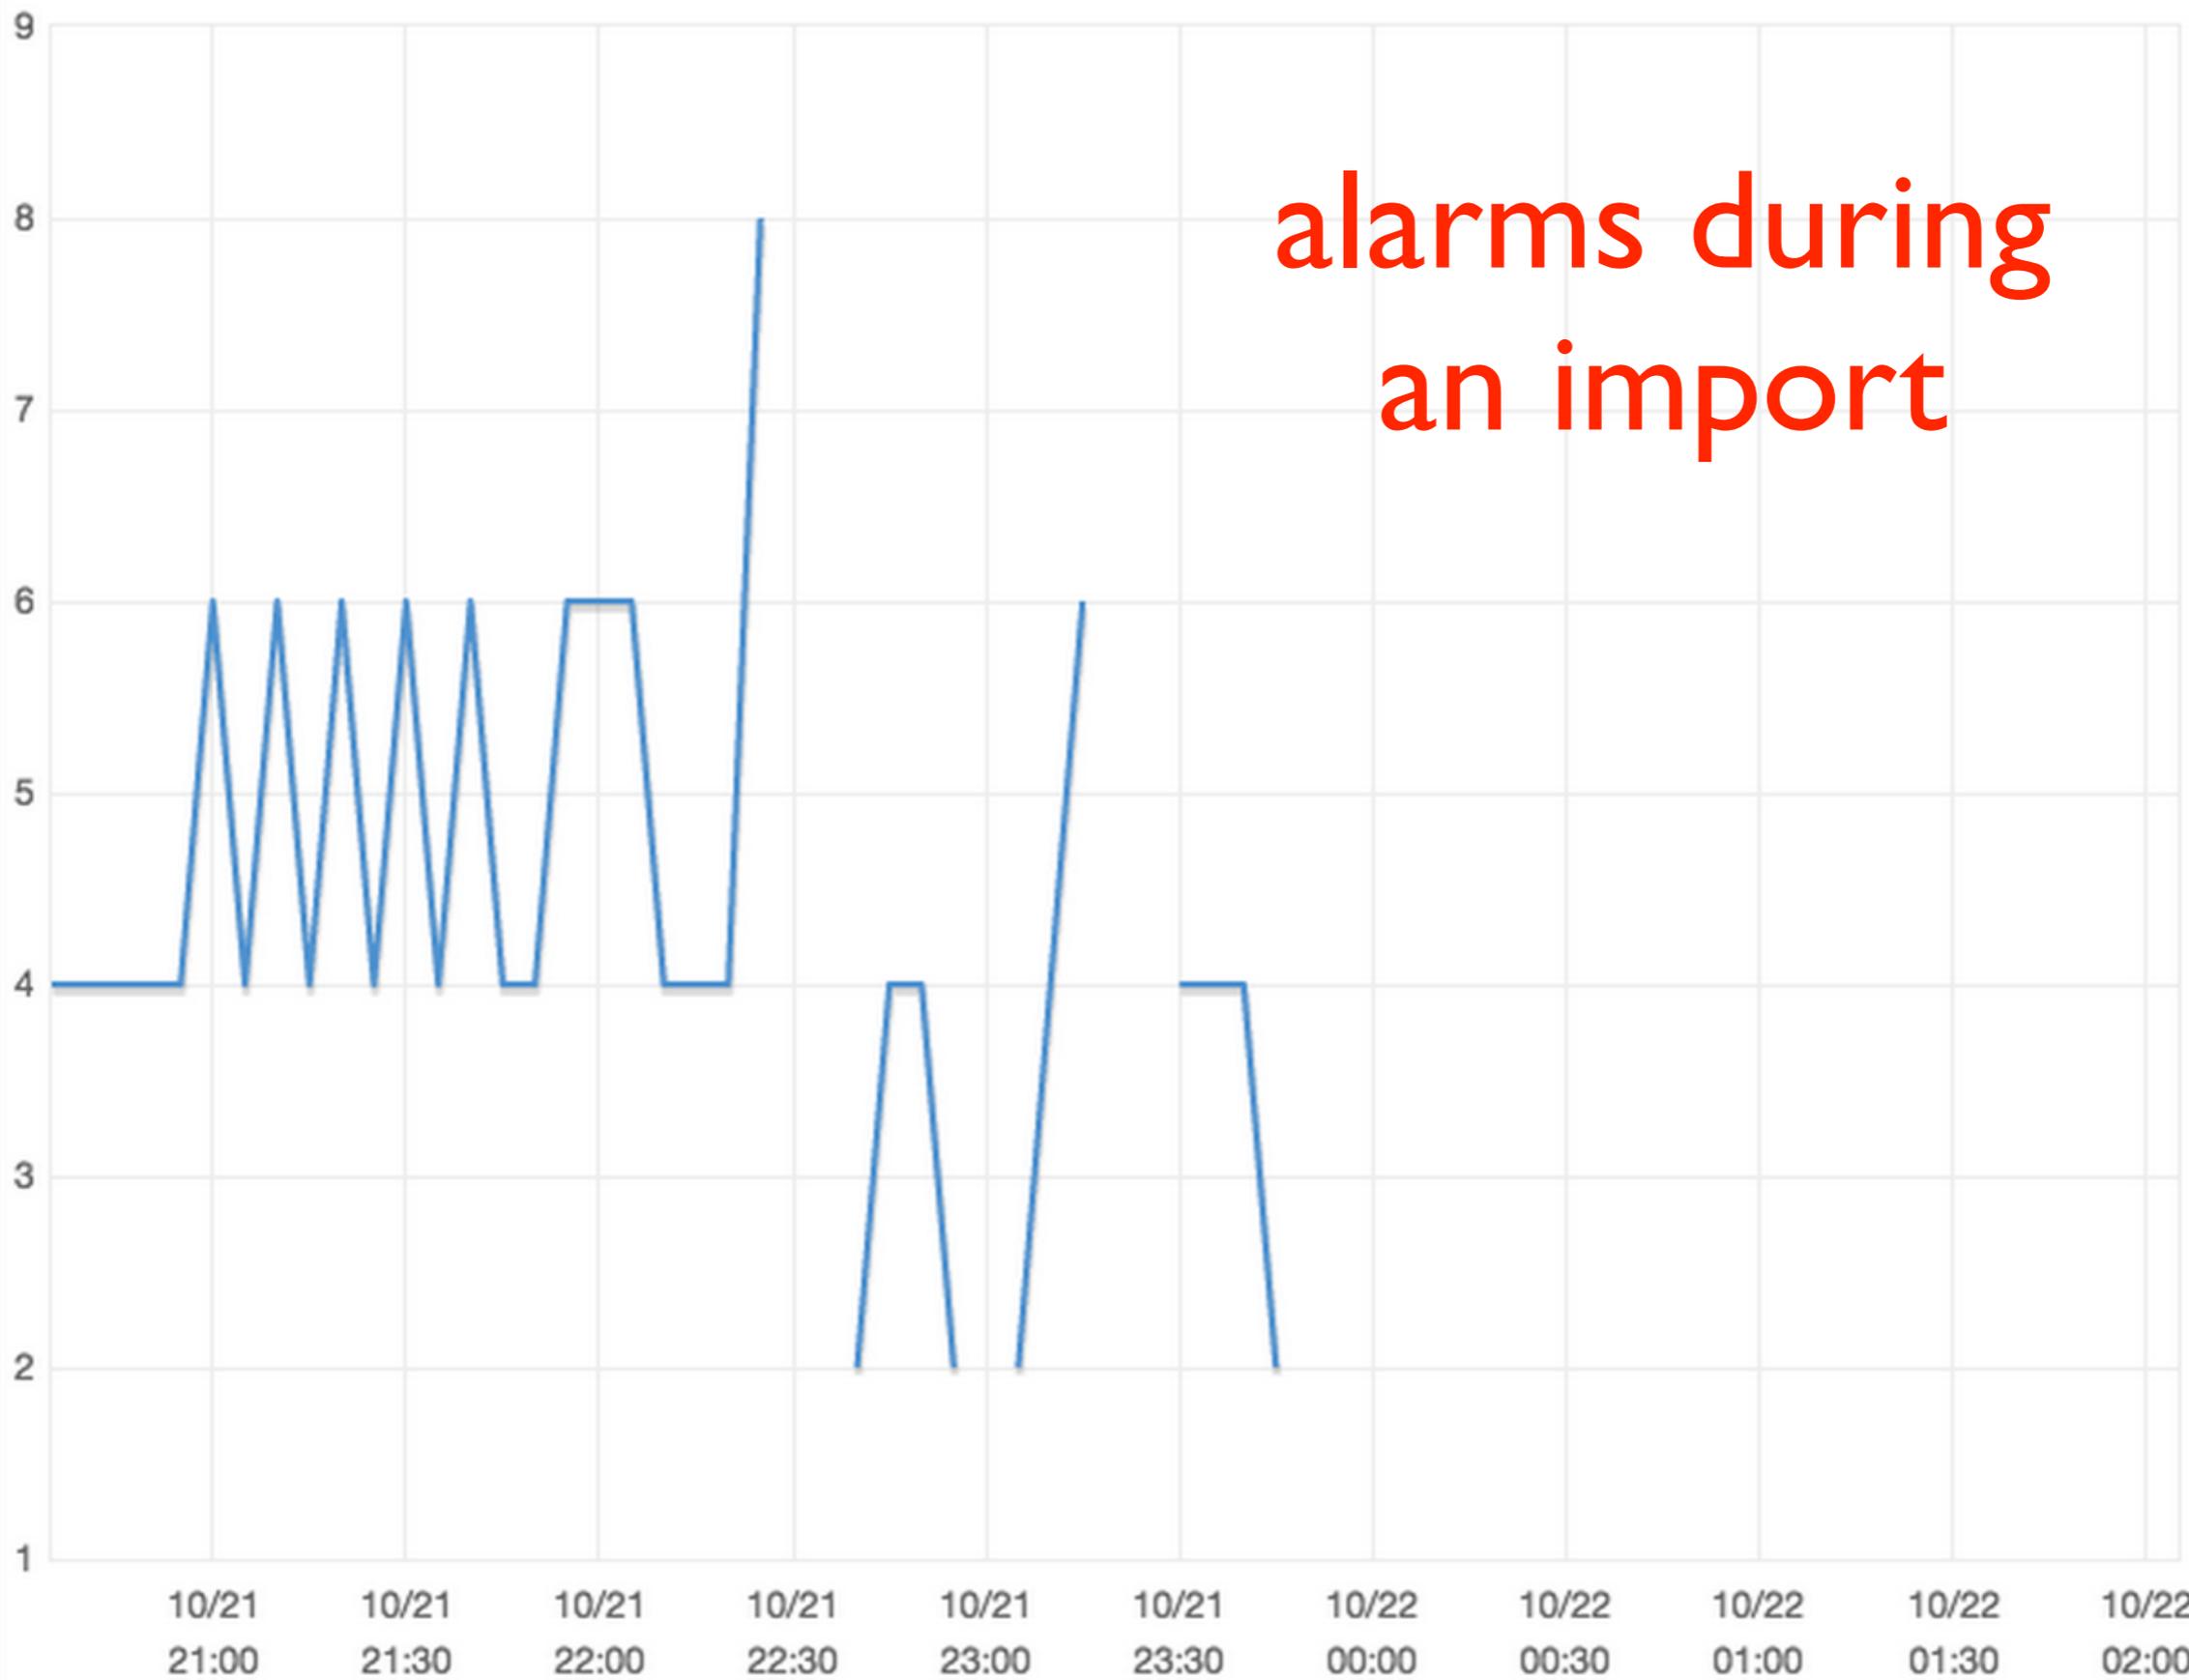
healthy transactor  
log writes



Left axis units: Milliseconds

LogWriteMsec





alarms during an import

Left axis units: Count

AlarmBackPressure



# capacity planning

data size	up to 10 billion datoms
write volume	consistent with data size in the long run
read volume	elastically scale peers as needed
housekeeping	<a href="http://docs.datomic.com/capacity.html#garbage-collection">http://docs.datomic.com/capacity.html#garbage-collection</a>
caching	test and monitor

# Datomic is a poor fit for

write scale

media storage (unstructured docs, audio, video)

churn (e.g. hit counter)

*big data, graphics, video, unstructured documents, log file analysis*

# Datomic is a good fit for

valuable information of record

dev and ops flexibility

history

read scale

*transactional data, business records, medical records, financial records, scientific records, inventory, configuration, web apps, departmental databases, cloud apps*

# lab: assess Datomic

for your dataset, create a comparison matrix

which Datomic characteristics are

beneficial?

detrimental?

irrelevant?

compared to ?

10 billion datoms  
access pattern at read-time  
ACID  
as-of  
attributes  
background indexing  
cardinality  
cloudwatch monitoring  
column access (AEVT)  
component attributes  
cross-db joins  
datalog  
data API  
document access  
edn

elastically scalable read  
entity maps  
excision  
graph access (VAET)  
high availability  
history  
joins  
lazy entities  
lookup refs  
memory speeds  
partitions  
pluggable monitoring  
pluggable storage  
pull  
query functions  
query predicates

raw index access  
reified transactions  
row access (EAVT)  
rules  
since  
sync  
time points  
transaction functions  
transparent memcached  
tx reports  
unique ids  
universal schema  
value access (AVET)  
upsert  
with

# thanks!

**@stuarthalloway**

<https://github.com/stuarthalloway/presentations/wiki>

<http://www.linkedin.com/pub/stu-halloway/0/110/543/>

<https://twitter.com/stuarthalloway>

<mailto:stu@cognitect.com>